



Linux debugging, profiling and tracing training

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Latest update: October 26, 2025.

Document updates and training details:
<https://bootlin.com/training/debugging>


Corrections, suggestions, contributions and translations are welcome!
Send them to feedback@bootlin.com





Linux debugging, profiling and tracing training

- ▶ These slides are the training materials for Bootlin's *Linux debugging, profiling and tracing* training course.
- ▶ If you are interested in following this course with an experienced Bootlin trainer, we offer:
 - **Public online sessions**, opened to individual registration. Dates announced on our site, registration directly online.
 - **Dedicated online sessions**, organized for a team of engineers from the same company at a date/time chosen by our customer.
 - **Dedicated on-site sessions**, organized for a team of engineers from the same company, we send a Bootlin trainer on-site to deliver the training.
- ▶ Details and registrations:
<https://bootlin.com/training/debugging>
- ▶ Contact: training@bootlin.com




Icon by Eucalyp, Flaticon



About Bootlin

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Bootlin introduction

- ▶ Engineering company
 - In business since 2004
 - Before 2018: *Free Electrons*
- ▶ Team based in France and Italy
- ▶ Serving **customers worldwide**
- ▶ **Highly focused and recognized expertise**
 - Embedded Linux
 - Linux kernel
 - Embedded Linux build systems
- ▶ **Strong open-source** contributor
- ▶ Activities
 - **Engineering** services
 - **Training** courses
- ▶ <https://bootlin.com>

The logo for bootlin, featuring the word "bootlin" in a lowercase, sans-serif font. A small orange triangle is positioned above the dot of the letter 'i'.



Bootlin engineering services

Bootloader / firmware development

U-Boot, Barebox,
OP-TEE, TF-A, .../

Linux kernel porting and driver development

Linux BSP development, maintenance and upgrade

Embedded Linux build systems

Yocto, OpenEmbedded,
Buildroot, ...

Embedded Linux integration

Boot time, real-time,
security, multimedia,
networking

Open-source upstreaming

Get code integrated
in upstream
Linux, U-Boot, Yocto,
Buildroot, ...



Bootlin training courses

Embedded Linux system development

On-site: 4 or 5 days
Online: 7 * 4 hours

Linux kernel driver development

On-site: 5 days
Online: 7 * 4 hours

Yocto Project system development

On-site: 3 days
Online: 4 * 4 hours

Buildroot system development

On-site: 3 days
Online: 5 * 4 hours

Embedded Linux networking

On-site: 3 days
Online: 4 * 4 hours

Understanding the Linux graphics stack

On-site: 2 days
Online: 4 * 4 hours

Embedded Linux audio

On-site: 2 days
Online: 4 * 4 hours

Real-Time Linux with PREEMPT_RT

On-site: 2 days
Online: 3 * 4 hours

Linux debugging, tracing, profiling and performance analysis

On-site: 3 days
Online: 4 * 4 hours

All our training materials are freely available
under a free documentation license (CC-BY-SA 3.0)
See <https://bootlin.com/training/>



Bootlin, an open-source contributor

- ▶ Strong contributor to the **Linux** kernel
 - In the top 30 of companies contributing to Linux worldwide
 - Contributions in most areas related to hardware support
 - Several engineers maintainers of subsystems/platforms
 - 9000 patches contributed
 - <https://bootlin.com/community/contributions/kernel-contributions/>
- ▶ Contributor to **Yocto Project**
 - Maintainer of the official documentation
 - Core participant to the QA effort
- ▶ Contributor to **Buildroot**
 - Co-maintainer
 - 6000 patches contributed
- ▶ Significant contributions to U-Boot, OP-TEE, Barebox, etc.
- ▶ Fully **open-source training materials**



Bootlin on-line resources

- ▶ Website with a technical blog:
<https://bootlin.com>
- ▶ Engineering services:
<https://bootlin.com/engineering>
- ▶ Training services:
<https://bootlin.com/training>
- ▶ LinkedIn:
<https://www.linkedin.com/company/bootlin>
- ▶ Elixir - browse Linux kernel sources on-line:
<https://elixir.bootlin.com>




Icon by Freepik, Flaticon



Generic course information

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Supported hardware

Discovery Kits from STMicroelectronics: STM32MP157A-DK1, STM32MP157D-DK1, STM32MP157C-DK2 or STM32MP157F-DK2

- ▶ STM32MP157 (Dual Cortex-A7 + Cortex-M4) CPU from STMicroelectronics
- ▶ 512 MB DDR3L RAM
- ▶ Gigabit Ethernet port
- ▶ 4 USB 2.0 host ports, 1 USB-C OTG port
- ▶ 1 Micro SD slot
- ▶ On-board ST-LINK/V2-1 debugger
- ▶ Misc: buttons, LEDs, audio codec
- ▶ LCD touchscreen (DK2 only)



DK1 Discovery Kit

Board and CPU documentation, design files, software: [A-DK1](#), [D-DK1](#), [C-DK2](#), [F-DK2](#)



Shopping list: hardware for this course

- ▶ STMicroelectronics STM32MP157D-DK1 Discovery kit
- ▶ USB-C cable for the power supply
- ▶ USB-A to micro B cable for the serial console
- ▶ RJ45 cable for networking
- ▶ A micro SD card with at least 128 MB of capacity





Training quiz and certificate

- ▶ You have been given a quiz to test your knowledge on the topics covered by the course. That's not too late to take it if you haven't done it yet!
- ▶ At the end of the course, we will submit this quiz to you again. That time, you will see the correct answers.
- ▶ It allows Bootlin to assess your progress thanks to the course. That's also a kind of challenge, to look for clues throughout the lectures and labs / demos, as all the answers are in the course!
- ▶ Another reason is that we only give training certificates to people who achieve at least a 50% score in the final quiz **and** who attended all the sessions.



Participate!

During the lectures...

- ▶ Don't hesitate to ask questions. Other people in the audience may have similar questions too.
- ▶ Don't hesitate to share your experience too, for example to compare Linux with other operating systems you know.
- ▶ Your point of view is most valuable, because it can be similar to your colleagues' and different from the trainer's.
- ▶ In on-line sessions
 - Please always keep your camera on!
 - Also make sure your name is properly filled.
 - You can also use the "Raise your hand" button when you wish to ask a question but don't want to interrupt.
- ▶ All this helps the trainer to engage with participants, see when something needs clarifying and make the session more interactive, enjoyable and useful for everyone.



Collaborate!

As in the Free Software and Open Source community, collaboration between participants is valuable in this training session:

- ▶ Use the dedicated Matrix channel for this session to add questions.
- ▶ If your session offers practical labs, you can also report issues, share screenshots and command output there.
- ▶ Don't hesitate to share your own answers and to help others especially when the trainer is unavailable.
- ▶ The Matrix channel is also a good place to ask questions outside of training hours, and after the course is over.

embedded-linux-nov2020 Channel for

Srinath

Michaelo: What should be CROSS_COMPILE variable set to in case of the Xplained board? I ran into some issues with my USB hub so doing the u-boot again

Michaelo

Srinath: you should look at the name of the cross-compiler in the toolchain's bin/ directory. CROSS_COMPILE should be set to what's before "gcc" in the name, including the trailing "-". Like if the compiler is arm-buildroot-linux-gcc, CROSS_COMPILE should be arm-buildroot-linux-

2 messages deleted.

Srinath

Will ask them here since I am going to do labs after the session is over! Thank!

Srinath changed their display name to srinath.

@srujanika-matrix.org

I tried to finalize Kernel - Cross-compiling task, but my system is not able to restart the new kernel. Does anyone know what can be the root cause?

Decrypt image.png (109.2 KB)

arnaud.a

I had the same because I accidentally renamed the vmlinux file from the kernel

Send an encrypted message...



Practical lab - Training Setup




Prepare your lab environment

- ▶ Download and extract the lab archive



Debugging, Tracing, Profiling


© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Debugging, Tracing, Profiling

- ▶ Debugging, tracing and profiling are often used for development purposes
- ▶ All of these methods have different goals which aim at perfecting the software that is being developed
- ▶ Requires some knowledge about underlying mechanisms to correctly identify and fix bugs





Debugging

- ▶ Finding and fixing bugs that might exist in your software/system
- ▶ Use of various tools and methods to achieve that
 - Interactive debugging (With GDB for instance)
 - Postmortem analysis (Using coredump for instance)
 - Control flow analysis (With tracing tools)
 - Testing (Targeted tests)
- ▶ Most commonly done through debuggers in development environment
- ▶ Generally intrusive, allowing to pause and resume execution

"Everyone knows that debugging is twice as hard as writing a program in the first place. So if you're as clever as you can be when you write it, how will you ever debug it?"

- Brian Kernighan



Tracing

- ▶ Following the execution flow of an application to understand the bottlenecks and problems.
- ▶ Achieved by instrumenting code either at compile time or runtime.
 - Can be done using specific tracers such as *LTtng*, *trace-cmd*, *SystemTap* etc
- ▶ Goes from the user space called functions up to the kernel ones
- ▶ Allows to identify functions and values that are used while application executes
- ▶ Often works by recording traces during runtime and then visualizing data.
 - Implies a large amount of recorded data since the complete execution trace is recorded
 - Often bigger overhead than profiling.
- ▶ Can also be used for debugging purpose since data can be extracted with tracepoints.



- ▶ Analysis at program runtime to assist performance optimizations
- ▶ Often achieved by sampling counters during execution
- ▶ Uses specific tools, libraries and operating system features to measure performance.
 - Using *perf*, *OProfile* for instance.
- ▶ First step consists in gathering data from program execution
 - Function call count, memory usage, CPU load, cache miss, etc
- ▶ Then extracting meaningful information from these data and modify the program to optimize it



Event sources

- ▶ Some activities like tracing or profiling involve some preliminary data collection
- ▶ Those data are gathered from different **event sources**, which can be of various types.
 - Some low-level events are gathered directly by the hardware (eg: CPU cycles, MMU exceptions...)
 - Some events are generated by some code explicitly added to generate traces, either in an application or in the kernel: those are **static tracepoints**
 - Some are generated by instrumentation added at runtime (ie without having to modify/rebuild the application and/or the kernel): those are **dynamic probes**




Linux Application Stack

© Copyright 2004-2025, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!





User/Kernel mode



- ▶ User mode vs Kernel mode are often used to refer to the privilege level of execution.
- ▶ This mode actually refers to the processor execution mode which is a hardware mode.
 - Might be named differently between architectures but the goal is the same
- ▶ Allows the kernel to control the full processor state (handle exceptions, MMU, etc) whereas the userspace can only do basic control and execute under the kernel supervision.



Introduction to Processes and Threads



Processes and Threads (1/2)

- ▶ A process is a group of resources that are allocated by the kernel to allow the execution of a program.
 - Memory regions, threads, file descriptors, etc.
- ▶ A process is identified by a PID (**P**rocess **ID**) and all the information that are specific to this process are exposed in `/proc/<pid>`.
 - A special file named `/proc/self` accessible by the process points to the proc folder associated to it.
- ▶ When starting a process, it initially has one execution thread that is represented by a `struct task_struct` and that can be scheduled.
 - A process is represented in the kernel by a thread associated to multiple resources.



Processes and Threads (2/2)

- ▶ Threads are independent execution units that are sharing common resources inside a process.
 - Same address space, file descriptors, etc.
- ▶ A new process is created using the `fork()` system call ([man 2 fork](#)) and a new thread is created using `pthread_create()` ([man 3 pthread_create](#)).
 - Internally, both will call `clone()` with different flags
- ▶ At any moment, only one task is executing on a CPU core and is accessible using [get_current\(\)](#) function (defined by architecture and often stored in a register).
- ▶ Each CPU core will execute a different task.
- ▶ A task can only be executing on one core at a time.



MMU and memory management




MMU and memory management

- ▶ Under Linux Kernel (when using `CONFIG_MMU=y`), all addresses that are accessed by the CPU are virtual
- ▶ The Memory Management Unit allows to map these virtual addresses to physical memory (either RAM or IO)
- ▶ All these mappings are inserted into the page table that is used by the MMU hardware to translate the CPU access from virtual to physical addresses
- ▶ The MMU allows to restrict access to the page mappings via some attributes
 - No Execute, Writable, Readable bits, Privileged/User bit, cacheability
- ▶ The MMU base unit for mappings is called a page
- ▶ Page size is fixed and depends on the architecture/kernel configuration.



Userspace/Kernel memory layout


- ▶ Each process has its own set of virtual memory areas (mm field of `struct task_struct`).
- ▶ Also have their own page table
 - But share the same kernel mappings
- ▶ By default, all user mapping addresses are randomized to minimize attack surface (base of heap, stack, text, data, etc).
 - **A**ddress **S**pace **L**ayout **R**andomization
 - Can be disabled using `norandmaps` command line parameter





Userspace/Kernel memory layout


Multiple processes have different user memory spaces





Kernel memory map

- ▶ The kernel has its own memory mapping.
- ▶ Linear mapping is setup at kernel startup by inserting all the entries in the kernel init page table.
- ▶ Multiple areas are identified and their location differs between the architectures.
- ▶ **Kernel Address Space Layout Randomization** also allows to randomize kernel address space layout.
 - Can be disabled using `nokaslr` command line parameter





Userspace memory segments

- ▶ When starting a process, the kernel sets up several *Virtual Memory Areas* (VMA), backed by `struct vm_area_struct`, with different execution attributes.
- ▶ VMA are actually memory zones that are mapped with specific attributes (R/W/X).
- ▶ A segmentation fault happens when a program tries to access an unmapped area or a mapped area with an access mode that is not allowed.
 - Writing data in a read-only segment
 - Executing data from a non-executable segment
- ▶ New memory zones can be created using `mmap()` ([man 2 mmap](#))
- ▶ Per application mappings are visible in `/proc/<pid>/maps`

```
7f1855b2a000-7f1855b2c000 rw-p 00030000 103:01 3408650 ld-2.33.so
7ffc01625000-7ffc01646000 rw-p 00000000 00:00 0 [stack]
7ffc016e5000-7ffc016e9000 r--p 00000000 00:00 0 [vvar]
7ffc016e9000-7ffc016eb000 r-xp 00000000 00:00 0 [vdso]
```



Userspace memory types

	Private	Shared
Anonymous	- stack - malloc() - brk()/sbrk() - mmap(PRIVATE, ANON)	- POSIX shm* - mmap(SHARED, ANON)
File-backed	- mmap(PRIVATE, fd) - program (code) - shared libraries	- mmap(SHARED, fd)



Terms for memory in Linux tools

- ▶ When using Linux tools, four terms are used to describe memory:
 - *VSS/VSZ*: Virtual Set Size (Virtual memory size, shared libraries included).
 - *RSS*: Resident Set Size (Total physical memory usage, shared libraries included).
 - *PSS*: Proportional Set Size (Actual physical memory used, divided by the number of times it has been mapped).
 - *USS*: Unique Set Size (Physical memory occupied by the process, shared mappings memory excluded).
- ▶ $VSS \geq RSS \geq PSS \geq USS$.




The process context



Process context

- ▶ The *process context* can be seen as the content of the CPU registers associated to a process: execution register, stack register...
- ▶ This context also designates an execution state and allows to sleep inside kernel mode.
- ▶ A process that is executing in process context can be preempted.
- ▶ While executing in such context, the current process `struct task_struct` can be accessed using `get_current()`.





Scheduling



Scheduling

- ▶ The scheduler can be invoked for various reasons
 - On a periodic tick caused by interrupt (`HZ`)
 - On a programmed interrupt on tickless systems (`CONFIG_NO_HZ=y`)
 - Voluntarily by calling `schedule()` in code
 - Implicitly by calling functions that can sleep (blocking operations such as `kmalloc()`, `wait_event()`).
- ▶ When entering the schedule function, the scheduler will elect a new `struct task_struct` to run and will eventually call the `switch_to()` macro.
- ▶ `switch_to()` is defined by architecture code and it will save the current task process context and restore the one of the next task to be run while setting the new current task running.



The Linux Kernel Scheduler

- ▶ The Linux Kernel Scheduler is a key piece in having a real-time behaviour
- ▶ It is in charge of deciding which **runnable** task gets executed
- ▶ It also elects on which CPU the task runs, and is tightly coupled to CPUidle and CPUFreq
- ▶ It schedules both **userspace** tasks and **kernel** tasks
- ▶ Each task is assigned one **scheduling class** or **policy**
- ▶ The class determines the algorithm used to elect each task
- ▶ Tasks with different scheduling classes can coexist on the system



Non-Realtime Scheduling Classes

There are 3 **Non-RealTime** classes

- ▶ **SCHED_OTHER**: The default policy, using a time-sharing algorithm
 - This policy is actually called **SCHED_NORMAL** by the kernel
- ▶ **SCHED_BATCH**: Similar to **SCHED_OTHER**, but designed for CPU-intensive loads that affect the wakeup time
- ▶ **SCHED_IDLE**: Very low priority class. Tasks with this policy will run only if nothing else needs to run.
- ▶ **SCHED_OTHER** and **SCHED_BATCH** use the **nice** value to increase or decrease their scheduling frequency
 - A higher nice value means that the tasks gets scheduled **less** often



Realtime Scheduling Classes

There are 3 **Realtime** classes

- ▶ Runnable tasks will preempt any other lower-priority task
- ▶ `SCHED_FIFO`: All tasks with the same priority are scheduled **First in, First out**
- ▶ `SCHED_RR`: Similar to `SCHED_FIFO` but with a time-sharing round-robin between tasks with the same priority
- ▶ Both `SCHED_FIFO` and `SCHED_RR` can be assigned a priority between 1 and 99
- ▶ `SCHED_DEADLINE`: For tasks doing recurrent jobs, extra attributes are attached to a task
 - A computation time, which represents the time the task needs to complete a job
 - A deadline, which is the maximum allowable time to compute the job
 - A period, during which only one job can occur
- ▶ Using one of these classes is necessary but not sufficient to get real-time behavior



Changing the Scheduling Class

- ▶ The Scheduling Class is set per-task, and defaults to `SCHED_OTHER`
- ▶ The `man 2 sched_setscheduler` syscall allows changing the class of a task
- ▶ The `chrt` tool uses it to allow changing the class of a running task:
 - `chrt -f/-b/-o/-r/-d -p PRIO PID`
- ▶ It can also be used to launch a new program with a dedicated class:
 - `chrt -f/-b/-o/-r/-d PRIO CMD`
- ▶ To show the current class and priority:
 - `chrt -p PID`
- ▶ New processes will inherit the class of their parent except if the `SCHED_RESET_ON_FORK` flag is set with `man 2 sched_setscheduler`
- ▶ See `man 7 sched` for more information



Context switching



Context switching

- ▶ Context switching is the action of changing the execution mode of the processor (Kernel $\leftrightarrow$ User).
 - Explicitly by executing system calls instructions (synchronous request to the kernel from user mode).
 - Implicitly when receiving exceptions (MMU fault, interrupts, breakpoints, etc).
- ▶ This state change will end up in a kernel entrypoint (often call vectors) that will execute necessary code to setup a correct state for kernel mode execution.
- ▶ The kernel takes care of saving registers, switching to the kernel stack and potentially other things depending on the architecture.
 - Does not use the user stack but a specific kernel fixed size stack for security purposes.



Exceptions

- ▶ Exceptions designate the kind of events that will trigger a CPU execution mode change to handle the exception.
- ▶ Two main types of exceptions exist: synchronous and asynchronous.
 - Asynchronous exceptions when a fault happens while executing (MMU, bus abort, etc) or when an interrupt is received (either software or hardware).
 - Synchronous when executing some specific instructions (breakpoint, syscall, etc)
- ▶ When such exception is triggered, the processor will jump to the exception vector and execute the code that was setup for this exception.




Interrupts

- ▶ Interrupts are asynchronous signals that are generated by the hardware peripherals.
 - Can also be synchronous when generated using a specific instruction (**I**nter **P**rocessor **I**nterrupts for instance).
- ▶ When receiving an interrupt, the CPU will change its execution mode by jumping to a specific vector and switching to kernel mode to handle the interrupt.
- ▶ When multiple CPUs (cores) are present, interrupts are often directed to a single core.
- ▶ This is called "IRQ affinity" and it allows to control the IRQ load for each CPU
 - See [core-api/irq/irq-affinity](#) and [man irqbalance\(1\)](#)



Interrupts

- ▶ While handling the interrupts, the kernel is executing in a specific context named *interrupt context*.
- ▶ This context does not have access to userspace and should not use `get_current()`.
- ▶ Depending on the architecture, might use an IRQ stack.
- ▶ Interrupts are disabled (no nested interrupt support)!





System Calls (1/2)

- ▶ A system call allows the user space to request services from the kernel by executing a special instruction that will switch to the kernel mode ([man 2 syscall](#))
 - When executing functions provided by the libc (`read()`, `write()`, etc), they often end up executing a system call.
- ▶ System calls are identified by a numeric identifier that is passed via the registers.
 - The kernel exports some defines (in `unistd.h`) that are named `__NR_<syscall>` and defines the syscall identifiers.

```
#define __NR_read 63  
#define __NR_write 64
```



System Calls (2/2)

- ▶ The kernel holds a table of function pointers which matches these identifiers and will invoke the correct handler after checking the validity of the syscall.
- ▶ System call parameters are passed via registers (up to 6).
- ▶ When executing this instruction the CPU will change its execution state and switch to the kernel mode.
- ▶ Each architecture uses a specific hardware mechanism ([man 2 syscall](#))

```
mov w8, #__NR_getpid
svc #0
tstne x0, x1
```



Kernel execution contexts



Kernel execution contexts

- ▶ The kernel runs code in various contexts depending on the event it is handling.
- ▶ Might have interrupts disabled, specific stack, etc.



Kernel threads

- ▶ Kernel threads (kthreads) are a special kind of `struct task_struct` that do not have any user resources associated (`mm == NULL`).
- ▶ These processes are cloned from the `kthreadd` process and can be created using `kthread_create()`.
- ▶ Kernel threads are scheduled and are allowed to sleep much like a process executing in process context.
- ▶ Kernel threads are visible and their names are displayed between brackets under `ps`:

```
$ ps --ppid 2 -p 2 -o uname,pid,ppid,cmd,cls
USER      PID    PPID  CMD                      CLS
root        2         0 [kthreadd]                TS
root        3         2 [rcu_gp]                   TS
root        4         2 [rcu_par_gp]               TS
root        5         2 [netns]                    TS
root        7         2 [kworker/0:0H-events_highpr TS
root       10         2 [mm_percpu_wq]             TS
root       11         2 [rcu_tasks_kthread]        TS
```



Workqueues


- ▶ Workqueues allows to schedule some work to be executed at some point in the future
- ▶ Workqueues are executing the work functions in kernel threads.
 - Allows to sleep while executing the deferred work.
 - Interrupts are enabled while executing
- ▶ Work can be executed either in dedicated work queues or in the default workqueue that is shared by multiple users.



- ▶ SoftIRQs is a specific kernel mechanism that is executed in software interrupt context.
- ▶ Allows to execute code that needs to be deferred after interrupt handling but needs low latency.
 - Executed right after hardware IRQ have been handled in interrupt context.
 - Same context as executing interrupt handler so sleeping is not allowed.
- ▶ Anyone wanting to run some code in softirq context should likely not create its own but prefer some entities implemented on top of it. There are for example tasklets, and the BH workqueues (Bottom Half workqueues) which aim to replace tasklets since 6.9.



Interrupts & Softirqs





Threaded interrupts

- ▶ Threaded interrupts are a mechanism that allows to handle the interrupt using a hard IRQ handler and a threaded IRQ handler.
 - Created calling `request_threaded_irq()` instead of `request_irq()`
- ▶ A threaded IRQ handler will allow to execute work that can potentially sleep in a kthread.
- ▶ One kthread is created for each interrupt line that was requested as a threaded IRQ.
 - *kthread* is named `irq/<irq>-<name>` and can be seen using *ps*.



Allocations and context

- ▶ Allocating memory in the kernel can be done using multiple functions:
 - `void *kmalloc(size_t size, gfp_t gfp_mask);`
 - `void *kzalloc(size_t size, gfp_t gfp_mask);`
 - `unsigned long __get_free_pages(gfp_t gfp_mask, unsigned int order)`
- ▶ All allocation functions take a `gfp_mask` parameter which allows to designate the kind of memory that is needed.
 - `GFP_KERNEL`: Normal allocation, can sleep while allocating memory (can not be used in interrupt context).
 - `GFP_ATOMIC`: Atomic allocation, won't sleep while allocating data.



Practical lab - Preparing the system




Prepare the STM32MP157D board

- ▶ Build an image using Buildroot
- ▶ Connect the board
- ▶ Load the kernel from SD card
- ▶ Mount the root filesystem over NFS



Linux Common Analysis & Observability Tools

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Pseudo Filesystems



Pseudo Filesystems

- ▶ Some virtual filesystems are exposed by the kernel and provide a lot of information on the system.
- ▶ *procfs* contains information about processes and system information.
 - Mounted on `/proc`
 - Often parsed by tools to display raw data in a more user-friendly way.
- ▶ *sysfs* provides information about hardware/logical devices, association between devices and drivers.
 - Mounted on `/sys`
- ▶ *debugfs* exposes information related to debug.
 - Typically mounted on `/sys/kernel/debug/`
 - `mount -t debugfs none /sys/kernel/debug`



- ▶ *procfs* exposes information about processes and system ([man 5 proc](#)).
 - `/proc/cpuinfo` CPU information.
 - `/proc/meminfo` memory information (used, free, total, etc).
 - `/proc/sys/` contains system parameters that can be tuned. The list of parameters that can be modified is available at [admin-guide/sysctl/index](#)
 - `/proc/interrupts`: interrupt count per CPU for each interrupt in use
 - We also have one entry per interrupt in `/proc/irq` for specific configuration/status for each interrupt line
 - `/proc/<pid>/` process related information
 - `/proc/<pid>/status` process basic information
 - `/proc/<pid>/maps` process memory mappings
 - `/proc/<pid>/fd` file descriptors of the process
 - `/proc/<pid>/task` descriptors of threads belonging to the process
 - `/proc/self/` will refer to the process used to access the file
- ▶ A list of all available *procfs* file and their content is described at [filesystems/proc](#) and [man 5 proc](#)



- ▶ `sysfs` filesystem exposes information about various kernel subsystems, hardware devices and association with drivers ([man 5 sysfs](#)).
- ▶ This allows to find the link between drivers and devices through a file hierarchy representing the kernel internal tree of devices.
- ▶ `/sys/kernel` contains interesting files for kernel debugging:
 - `irq` with information about interrupts (mapping, count, etc).
 - `tracing` for tracing control.
- ▶ [admin-guide/abi-stable](#)



debugfs

- ▶ *debugfs* is a simple RAM-based filesystem which exposes debugging information.
- ▶ Used by some subsystems (*clk*, *block*, *dma*, *gpio*, etc) to expose debugging information related to the internals.
- ▶ Usually mounted on `/sys/kernel/debug`
 - Dynamic debug features exposed through `/sys/kernel/debug/dynamic_debug` (also exposed in `proc`)
 - Clock tree exposed through `/sys/kernel/debug/clk/clk_summary`.




ELF file analysis



ELF files

Executable and Linkable Format

- ▶ File starting with a header which holds binary structures defining the file
- ▶ Collection of segments and sections that contain data
 - `.text` section: Code
 - `.data` section: Data
 - `.rodata` section: Read-only Data
 - `.debug_info` section: Contains debugging information
- ▶ Sections are part of a segment which can be loadable in memory
- ▶ Same format for all architectures supported by the kernel and also `vmlinux` format
 - Also used by a lot of other operating systems as the standard executable file format





binutils for ELF analysis

- ▶ The binutils are used to deal with binary files, either object files or executables.
 - Includes `ld`, `as` and other useful tools.
- ▶ *readelf* displays information about ELF files (header, section, segments, etc).
- ▶ *objdump* allows to display information and disassemble ELF files.
- ▶ *objcopy* can convert ELF files or extract/translate some parts of it.
- ▶ *nm* displays the list of symbols embedded in ELF files.
- ▶ *addr2line* finds the source code line/file pair from an address using an ELF file with debug information



binutils example (1/2)

- ▶ Finding the address of `ksys_read()` kernel function using *nm*:

```
$ nm vmlinux | grep ksys_read
c02c7040 T ksys_read
```

- ▶ Using *addr2line* to match a kernel OOPS address or a symbol name with source code:

```
$ addr2line -s -f -e vmlinux ffffffff8145a8b0
queue_wc_show
blk-sysfs.c:516
```



binutils example (2/2)

- ▶ Display an ELF header with *readelf*:

```
$ readelf -h binary
ELF Header:
Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
Class:                               ELF64
Data:                               2's complement, little endian
Version:                             1 (current)
OS/ABI:                             UNIX - System V
ABI Version:                         0
Type:                               DYN (Position-Independent Executable file)
Machine:                             Advanced Micro Devices X86-64
...
```

- ▶ Convert an ELF file to a flat binary file using *objcopy*:

```
$ objcopy -O binary file.elf file.bin
```



- ▶ In order to display the shared libraries used by an ELF binary, one can use *ldd* (Generally packaged with C library. See [man 1 ldd](#)).
- ▶ *ldd* will list all the libraries that were used at link time.
 - Libraries that are loaded at runtime using `dlopen()` are not displayed.

```
$ ldd /usr/bin/bash
linux-vdso.so.1 (0x00007ffdf3fc6000)
libreadline.so.8 => /usr/lib/libreadline.so.8 (0x00007fa2d2aef000)
libc.so.6 => /usr/lib/libc.so.6 (0x00007fa2d2905000)
libncursesw.so.6 => /usr/lib/libncursesw.so.6 (0x00007fa2d288e000)
/lib64/ld-linux-x86-64.so.2 => /usr/lib64/ld-linux-x86-64.so.2 (0x00007fa2d2c88000)
```



Monitoring tools



Monitoring Tools

- ▶ Lots of monitoring tools on Linux to allow monitoring various part of the system.
- ▶ Most of the time, these are CLI interactive programs.
 - Processes with *ps*, *top*, *htop*, etc
 - Memory with *free*, *vmstat*
 - Networking
- ▶ Almost all these tools rely on the *sysfs* or *procfs* filesystem to obtain the processes, memory and system information but will display them in a more human-readable way.
 - Networking tools use a netlink interface with the networking subsystem of the kernel.



Process and CPU monitoring tools



Processes with *ps*

- ▶ The *ps* command allows to display a snapshot of active processes and their associated information ([man 1 ps](#))
 - Lists both user processes and kernel threads.
 - Displays PID, CPU usage, memory usage, uptime, etc.
 - Uses */proc/<pid>/* directory to obtain process information.
 - Almost always present on embedded platforms (provided by *Busybox*).
- ▶ By default, displays only the current user/current tty processes, but output is highly customizable:
 - *aux/-e*: show all processes
 - *-L*: show threads
 - *-p*: target a specific process
 - *-o*: select output columns to display
- ▶ Useful for scripting and parsing since its output is static.



Processes with *ps*

- Display all processes in a friendly way:

```
$ ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	168864	12800	?	Ss	09:08	0:00	/sbin/init
root	2	0.0	0.0	0	0	?	S	09:08	0:00	[kthreadd]
root	3	0.0	0.0	0	0	?	I<	09:08	0:00	[rcu_gp]
root	4	0.0	0.0	0	0	?	I<	09:08	0:00	[rcu_par_gp]
root	5	0.0	0.0	0	0	?	I<	09:08	0:00	[netns]
...										
root	914	0.0	0.0	396216	16220	?	Ss1	09:08	0:04	/usr/libexec/udisks2/udisksd
avahi	929	0.0	0.0	8728	412	?	S	09:08	0:00	avahi-daemon: chroot helper
root	956	0.0	0.1	260304	19024	?	Ss1	09:08	0:02	/usr/sbin/NetworkManager --n
root	960	0.0	0.0	17040	5704	?	Ss	09:08	0:00	/sbin/wpa_supplicant -u -s -
root	962	0.0	0.0	317644	11896	?	Ss1	09:08	0:00	/usr/sbin/ModemManager
vnstat	987	0.0	0.0	5516	3696	?	Ss	09:08	0:00	/usr/sbin/vnstatd -n



Processes with *top*

- ▶ *top* command output information similar to *ps* but dynamic and interactive ([man 1 top](#)).
 - Also almost always present on embedded platforms (provided by *Busybox*)

```
$ top
top - 18:38:11 up 9:29, 1 user, load average: 2.84, 2.74, 2.02
Tasks: 371 total, 1 running, 370 sleeping, 0 stopped, 0 zombie
%Cpu(s): 5.8 us, 2.1 sy, 0.0 ni, 77.4 id, 14.7 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 15947.6 total, 1476.9 free, 7685.7 used, 6784.9 buff/cache
MiB Swap: 15259.0 total, 15238.7 free, 20.2 used. 7742.3 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
2988	cleger	20	0	5184816	1.2g	430244	S	26.7	7.9	60:24.27	firefox-esr
4326	cleger	20	0	16.4g	208104	81504	S	26.7	1.3	9:27.33	code
909	root	-51	0	0	0	0	S	13.3	0.0	15:12.15	irq/104-nvidia
41704	cleger	20	0	38.4g	373744	116984	S	13.3	2.3	13:25.76	code
91926	cleger	20	0	2514784	145360	95144	S	13.3	0.9	1:29.85	Web Content



- ▶ *mpstat* displays Multiprocessor statistics ([man 1 mpstat](#)).
- ▶ Useful to detect unbalanced CPU workloads, bad IRQ affinity, etc.

```
$ mpstat -P ALL
```

```
Linux 6.0.0-1-amd64 (fixe)
```

```
19/10/2022
```

```
_x86_64_
```

```
(4 CPU)
```

17:02:50	CPU	%usr	%nice	%sys	%iowait	%irq	%soft	%steal	%guest	%gnice	%idle
17:02:50	all	6,77	0,00	2,09	11,67	0,00	0,06	0,00	0,00	0,00	79,40
17:02:50	0	6,88	0,00	1,93	8,22	0,00	0,13	0,00	0,00	0,00	82,84
17:02:50	1	4,91	0,00	1,50	8,91	0,00	0,03	0,00	0,00	0,00	84,64
17:02:50	2	6,96	0,00	1,74	7,23	0,00	0,01	0,00	0,00	0,00	84,06
17:02:50	3	9,32	0,00	2,80	54,67	0,00	0,00	0,00	0,00	0,00	33,20
17:02:50	4	5,40	0,00	1,29	4,92	0,00	0,00	0,00	0,00	0,00	88,40



Memory monitoring tools



- ▶ *free* is a simple program that displays the amount of free and used memory in the system ([man 1 free](#)).
 - Useful to check if the system suffers from memory exhaustion
 - Uses `/proc/meminfo` to obtain memory information.

```
$ free -h
```

	total	used	free	shared	buff/cache	available
Mem:	15Gi	7.5Gi	1.4Gi	192Mi	6.6Gi	7.5Gi
Swap:	14Gi	20Mi	14Gi			

- ▶ *A small free value does not mean that your system suffers from memory depletion! Linux considers any unused memory as "wasted" so it uses it for buffers and caches to optimize performance. See also `drop_caches` from [man 5 proc](#) to observe buffers/cache impact on free/available memory*



- ▶ `vmstat` displays information about system virtual memory usage
- ▶ Can also display stats from processes, memory, paging, block IO, traps, disks and cpu activity ([man 8 vmstat](#)).
- ▶ Can be used to gather data at periodic interval using `vmstat <interval> <number>`

```
$ vmstat 1 6
```

procs		-----memory-----				---swap--		-----io----		-system--		-----cpu-----				
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
3	0	253440	1237236	194936	9286980	3	6	186	540	134	157	3	5	82	10	0

- ▶ *Note: `vmstat` consider a kernel block to be 1024 bytes*



- ▶ pmap displays process mappings more easily than accessing `/proc/<pid>/maps` ([man 1 pmap](#)).

```
# pmap 2002
2002: /usr/bin/dbus-daemon --session --address=systemd: --nofork --nopidfile --systemd-activation --syslog-only
...
00007f3f958bb000      56K r---- libdbus-1.so.3.32.1
00007f3f958c9000     192K r-x-- libdbus-1.so.3.32.1
00007f3f958f9000      84K r---- libdbus-1.so.3.32.1
00007f3f9590e000       8K r---- libdbus-1.so.3.32.1
00007f3f95910000       4K rw--- libdbus-1.so.3.32.1
00007f3f95937000       8K rw--- [ anon ]
00007f3f95939000       8K r---- ld-linux-x86-64.so.2
00007f3f9593b000     152K r-x-- ld-linux-x86-64.so.2
00007f3f95961000      44K r---- ld-linux-x86-64.so.2
00007f3f9596c000       8K r---- ld-linux-x86-64.so.2
00007f3f9596e000       8K rw--- ld-linux-x86-64.so.2
00007ffe13857000     132K rw--- [ stack ]
00007ffe13934000      16K r---- [ anon ]
00007ffe13938000       8K r-x-- [ anon ]
total                11088K
```



I/O monitoring tools



iostat

- ▶ *iostat* displays information about IOs per device on the system.
- ▶ Useful to see if a device is overloaded by IOs.

```
$ iostat
```

```
Linux 5.19.0-2-amd64 (fixe)      11/10/2022      _x86_64_      (12 CPU)
```

```
avg-cpu:  %user   %nice %system %iowait  %steal   %idle
           8,43    0,00    1,52    8,77    0,00   81,28
```

Device	tps	kB_read/s	kB_wrtn/s	kB_dscd/s	kB_read	kB_wrtn	kB_dscd
nvme0n1	55,89	1096,88	149,33	0,00	5117334	696668	0
sda	0,03	0,92	0,00	0,00	4308	0	0
sdb	104,42	274,55	2126,64	0,00	1280853	9921488	0



iotop

- ▶ *iotop* displays information about IOs much like *top* for each process.
- ▶ Useful to find applications generating too much I/O traffic.
 - Needs `CONFIG_TASKSTATS=y`, `CONFIG_TASK_DELAY_ACCT=y` and `CONFIG_TASK_IO_ACCOUNTING=y` to be enabled in the kernel.
 - Also needs to be enabled at runtime: `sysctl -w kernel.task_delayacct=1`

```
# iotop
Total DISK READ:      20.61 K/s | Total DISK WRITE:      51.52 K/s
Current DISK READ:    20.61 K/s | Current DISK WRITE:    24.04 K/s
   TID  PRIO  USER      DISK READ DISK WRITE>      COMMAND
   2629 be/4  cleger    20.61 K/s  44.65 K/s firefox-esr [Cache2 I/O]
   322  be/3  root       0.00 B/s   3.43 K/s [jbd2/nvme0n1p1-8]
 39055 be/4  cleger     0.00 B/s   3.43 K/s firefox-esr [DOMCacheThread]
     1  be/4  root       0.00 B/s   0.00 B/s init
     2  be/4  root       0.00 B/s   0.00 B/s [kthreadd]
     3  be/0  root       0.00 B/s   0.00 B/s [rcu_gp]
     4  be/0  root       0.00 B/s   0.00 B/s [rcu_par_gp]
...

```



Networking observability tools



- ▶ `ss` shows the status of network sockets
 - IPv4 and IPv6, UDP, TCP, ICMP and UNIX domain sockets
- ▶ Replaces *netstat*, now obsolete
- ▶ Gets info from `/proc/net`
- ▶ Usage:
 - `ss` by default shows connected sockets
 - `ss -l` shows listening sockets
 - `ss -a` shows both listening and connected sockets
 - `ss -4/-6/-x` shows only IPv4, IPv6, or UNIX sockets
 - `ss -t/-u` shows only TCP or UDP sockets
 - `ss -p` shows process using each socket
 - `ss -n` shows numeric addresses
 - `ss -s` shows a summary of existing sockets
- ▶ See [the ss manpage](#) for all the options



ss example output

```
# ss
Netid State  Recv-Q Send-Q           Local Address:Port           Peer Address:Port  Process
u_dgr ESTAB   0      0                * 304840                * 26673
u_str ESTAB   0      0      /run/dbus/system_bus_socket 42871                * 26100
icmp6 UNCONN  0      0                *:ipv6-icmp            *: *
udp   ESTAB   0      0      192.168.10.115%wlp0s20f3:bootpc 192.168.10.88:bootps
tcp   ESTAB   0      136           172.16.0.1:41376        172.16.11.42:ssh
tcp   ESTAB   0      273           192.168.1.77:55494      87.98.181.233:https
tcp   ESTAB   0      0      [2a02:....dbdc]:38466    [2001:....9]:imap2
...
#
```



- ▶ *iftop* displays bandwidth usage on an interface by remote host
- ▶ Visualizes bandwidth using histograms
- ▶ `iftop -i eth0`

	19.1Mb	38.1Mb	57.2Mb	76.3Mb	95.4Mb	
localhost.localnet		=> ams.source.kernel.org		424Kb	352Kb	403Kb
		<=		18.8Mb	15.6Mb	18.2Mb
localhost.localnet		=> bootlin.com		0b	113Kb	35.0Kb
		<=		0b	6.20Mb	1.84Mb
TX:		cum: 3.06MB	peak: 836Kb	rates: 424Kb 465Kb 439Kb		
RX:		146MB	44.3Mb	18.8Mb 21.8Mb 20.1Mb		
TOTAL:		149MB	45.1Mb	19.2Mb 22.3Mb 20.5Mb		

- ▶ The output can be customized interactively
- ▶ See [the iftop manpage](#) for details



- ▶ *tcpdump* allows to capture network traffic and decode many protocols
- ▶ `tcpdump -i eth0`
- ▶ based on the *libpcap* library for packet capture
- ▶ It can also store captured packets to a file and read them back
 - In the *pcap* format or the newer *pcapng* format
 - `tcpdump -i eth0 -w capture.pcap`
 - `tcpdump -r capture.pcap`
- ▶ A BPF capture filter can be used to avoid capturing irrelevant packets
 - `tcpdump -i eth0 tcp and not port 22`



tcpdump example output

```
# tcpdump -i eth0
18:41:22.913058 IP localhost.localnet.40764 > srv.localnet: 14324+ AAAA? bootlin.com. (29)
18:41:22.913797 IP srv.localnet > localhost.localnet.40764: 14324 0/1/0 (89)
18:41:22.914268 IP localhost.localnet > bootlin.com: ICMP echo request, id 3, seq 1, length 64
18:41:23.933063 IP localhost.localnet > bootlin.com: ICMP echo request, id 3, seq 2, length 64
18:41:24.957027 IP localhost.localnet > bootlin.com: ICMP echo request, id 3, seq 3, length 64
18:41:24.996415 IP bootlin.com > localhost.localnet: ICMP echo reply, id 3, seq 3, length 64
^C
# tcpdump -i eth0 tcp and not port 22
... IP B.https > A.38910: Flags [.] , ack 469, win 501, options [...], length 0
... IP B.https > A.38910: Flags [P.] , seq 2602:2857, ack 469, win 501, options [...], length 255
... IP A.38910 > B.https: Flags [.] , ack 2857, win 501, options [...], length 0
... IP A.38910 > B.https: Flags [P.] , seq 469:621, ack 2857, win 501, options [...], length 152
... IP B.https > A.38910: Flags [.] , ack 621, win 501, options [...], length 0
... IP B.https > A.38910: Flags [P.] , seq 2857:3825, ack 621, win 501, options [...], length 968
... IP A.38910 > B.https: Flags [P.] , seq 621:779, ack 3825, win 501, options [...], length 158
^C
#
```



- ▶ Similar to tcpdump, but with a GUI
- ▶ Also based on libpcap
 - Can capture and use the same BPF capture filters
 - Can load and save the same file formats
 - Useful for embedded: capture on the target with tcpdump, analyze on the host with Wireshark
- ▶ Has *dissectors* to decode hundreds of protocols
 - Each individual value from each packet is dissected into a separate field
 - Fields are very fine-grained, at least for the most common protocols
- ▶ Has *display filters* that allow filtering *already captured* packets
 - Each dissected field is also a filter key
- ▶ Can also capture and decode Bluetooth, USB, D-Bus and more



capture.pcap

File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

ip.addr eq 87.98.181.233 and ip.addr eq

No.	Time	Source	Destination	Protocol	Length	Sequence	Acknowledgment	Info
3	0.036445		87.98.181.233	TCP	66	1	1	42524 → 443 [ACK] Seq=1 Ack=1 Win=64256 Len=0 TSval=62317059 TSecr=62317059
4	0.037180		87.98.181.233	TLSv1.3	471	1	1	Client Hello
5	0.072674	87.98.181.233		TCP	66	1	406	443 → 42524 [ACK] Seq=1 Ack=406 Win=64768 Len=0 TSval=3932229694 TSecr=62317059
6	0.077240	87.98.181.233		TLSv1.3	1506	1	406	Server Hello, Change Cipher Spec, Application Data
7	0.077292		87.98.181.233	TCP	66	406	1441	42524 → 443 [ACK] Seq=406 Ack=1441 Win=64128 Len=0 TSval=623171059 TSecr=62317059

Frame 8: 1445 bytes on wire (11560 bits), 1445 bytes captured (11560 bits) on interface 0

Ethernet II, Src: AVMAudio (3c:a6:2f:12:34:56), Dst: PartIIR_54 (08:00:27:12:34:56)

Internet Protocol Version 4, Src: 87.98.181.233, Dst: 192.168.178.75

- 0100 = Version: 4
- 0101 = Header Length: 20 bytes (5)
- > Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
- Total Length: 1431
- Identification: 0xf881 (63617)
- Flags: 0x40, Don't fragment
 - 0... = Reserved bit: Not set
 - 1... = Don't fragment: Set
 - ...0 = More fragments: Not set
- ...0 0000 0000 0000 = Fragment Offset: 0
- Time to Live: 47
- Protocol: TCP (6)
- Header Checksum: 0xcd9f [validation disabled]
- [Header checksum status: Unverified]
- Source Address: 87.98.181.233
- Destination Address: 192.168.178.75

Transmission Control Protocol, Src Port: 443, Dst Port: 42524, Seq: 1441, Len: 0

- Source Port: 443
- Destination Port: 42524
- [Stream index: 0]
- [Conversation completeness: Complete, WITH_DATA (63)]
- [TCP Segment Len: 1379]
- Sequence Number: 1441 (relative conversation number)

Don't fragment (ip.flags.df, 1 byte(s))

Frame (1445 bytes) Reassembled TCP (2447 bytes)

Packets: 385 · Displayed: 385 (100.0%) Profile: Default




Check what is running on a system and its load

- ▶ Observe processes and IOs
- ▶ Display memory mappings
- ▶ Monitor resources



Application Debugging

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Good practices



Good practices

- ▶ Some good practices can allow you to save time before even needing to use a debugger
- ▶ Compiler are now smart enough to detect a wide range of errors at compile-time using warnings
 - Using `-Werror -Wall -Wextra` is recommended if possible to catch errors as early as possible
- ▶ Compilers now offer static analysis capabilities
 - GCC allows to do so using the `-fanalyzer` flag
 - LLVM provides **dedicated tools** that can be used in build process
- ▶ You can also enable component-specific helpers/hardening
 - If you are using the GNU C library, you can for example enable **`__FORTIFY_SOURCE`** macro to add runtime checks on inputs (e.g: buffers)



Building with debug information



Debugging with ELF files

- ▶ GDB uses ELF files since they contain the debugging information
- ▶ Debugging information uses the DWARF format
- ▶ Allows the debugger to match addresses and symbol names, call sites, etc
- ▶ Debugging information is generated by the compiler and included in the ELF file when compiled with `-g`
 - `-g1`: minimal debug information (enough for backtraces)
 - `-g2`: default debug level when using `-g`
 - `-g3`: includes extra debugging information (macro definitions)
- ▶ See [GCC documentation](#) about debugging for more information





Debugging with compiler optimizations

- ▶ Compiler optimizations (`-O<level>`) can lead to optimizing out some variables and function calls.
- ▶ Trying to display them with GDB will display
 - `$1 = <value optimized out>`
- ▶ If one wants to inspect variables and functions, it is possible to compile the code using `-O0` (no optimization).
 - *Note: The kernel can only be compiled with `-O2` or `-Os`*
- ▶ It is also possible to annotate function with compiler attributes:
 - `__attribute__((optimize("O0")))`
- ▶ Remove function `static` qualifier to avoid inlining the function
 - *Note: LTO (Link Time Optimization) can defeat this.*
- ▶ Set a specific variable as `volatile` to prevent the compiler from optimizing it out.



Instrumenting code crashes



Instrumenting code crashes

- ▶ Call stacks can be manually requested/generated in applications
- ▶ This can be done either with specialized libraries (eg: [libunwind](#)) or even directly the C library, e.g by using `backtrace()` ([man 3 backtrace](#)) with glibc:

```
char **backtrace_symbols(void *const *buffer, int size);
```

- ▶ We can wire those custom backtraces on signals thanks to `sigaction()` ([man 2 sigaction](#))
 - SIGSEGV signal can be hooked to dump a call stack of an invalid memory access

```
int sigaction(int sig, const struct sigaction *restrict act,  
              struct sigaction *restrict oact);
```

- ▶ Some libraries like [libsigsegv](#) ease the instrumentation of those signals



Custom code crash report (1/2)

```
[...]
void callee(void *ptr) {
    int *myptr = (int *)ptr;
    printf("Executing suspicious operation\n");
    myptr[2] = 0;
}

void caller(void) {
    void *ptr = NULL;
    callee(ptr);
}

void segfault_handler(int sig) {
    void *array[20];
    size_t size;
    char msg[] = "Segmentation fault!\n";

    write(STDERR_FILENO, msg, sizeof(msg));
    size = backtrace(array, 20);
    backtrace_symbols_fd(array, size, STDERR_FILENO);
    exit(1);
}
```

```
int main() {
    const struct sigaction act = {
        .sa_handler = segfault_handler,
        .sa_mask = 0,
        .sa_flags = 0,
    };
    if (sigaction(SIGSEGV, &act, NULL))
        exit(EXIT_FAILURE);
    printf("Calling a faulty function\n");
    caller();
    return 0;
}
```



Custom code crash report (2/2)

```
[root@arch-bootlin-alexis custom_backtrace]# ./main
Calling a faulty function
Executing suspicious operation
Segmentation fault!
./main(segfault_handler+0x60)[0x55c6e4c1723c]
/usr/lib/libc.so.6(+0x38f50)[0x7fecb0a95f50]
./main(callee+0x2b)[0x55c6e4c171b4]
./main(caller+0x1c)[0x55c6e4c171d9]
./main(main+0x2c)[0x55c6e4c1729a]
/usr/lib/libc.so.6(+0x23790)[0x7fecb0a80790]
/usr/lib/libc.so.6(__libc_start_main+0x8a)[0x7fecb0a8084a]
./main(_start+0x25)[0x55c6e4c170b5]
```

- ▶ Writing robust signal handlers is not easy
 - When your application has received a SIGSEGV signal, its execution can not really be trusted anymore
 - We are for example supposed to use only [reentrant functions](#) in signal handlers
 - Not following this rule may lead to undefined behavior



The ptrace system call



- ▶ The *ptrace* mechanism allows processes to trace other processes by accessing tracee memory and register contents
- ▶ A tracer can observe and control the execution state of another process
- ▶ Works by attaching to a tracee process using the `ptrace()` system call (see [man 2 ptrace](#))
- ▶ Can be executed directly using the `ptrace()` call but often used indirectly using other tools.

```
long ptrace(enum __ptrace_request request, pid_t pid, void *addr, void *data);
```

- ▶ Used by *GDB*, *strace* and all debugging tools that need access to the tracee process state



GDB



GDB: GNU Project Debugger

- ▶ The debugger on GNU/Linux, available for most embedded architectures.
- ▶ Supported languages: C, C++, Pascal, Objective-C, Fortran, Ada...
- ▶ Command-line interface
- ▶ Integration in many graphical IDEs
- ▶ Can be used to
 - control the execution of a running program, set breakpoints or change internal variables
 - to see what a program was doing when it crashed: post mortem analysis
- ▶ <https://www.gnu.org/software/gdb/>
- ▶ <https://en.wikipedia.org/wiki/Gdb>
- ▶ New alternative: *lldb* (<https://lldb.llvm.org/>) from the LLVM project.





GDB crash course (1/3)

- ▶ GDB is used mainly to debug a process by starting it with *gdb*
 - `$ gdb <program>`
- ▶ GDB can also be attached to running processes using the program PID
 - `$ gdb -p <pid>`
- ▶ When using GDB to start a program, the program needs to be run with
 - `(gdb) run [prog_arg1 [prog_arg2] ...]`



GDB crash course (2/3)

A few useful GDB commands

- ▶ `break foobar (b)`
Put a breakpoint at the entry of function `foobar()`
- ▶ `break foobar.c:42`
Put a breakpoint in `foobar.c`, line 42
- ▶ `print var`, `print $reg` or `print task->files[0].fd (p)`
Print the variable `var`, the register `$reg` or a more complicated reference. GDB can also nicely display structures with all their members
- ▶ `info registers`
Display architecture registers



GDB crash course (3/3)

- ▶ `continue (c)`
Continue the execution after a breakpoint
- ▶ `next (n)`
Continue to the next line, stepping over function calls
- ▶ `step (s)`
Continue to the next line, entering into subfunctions
- ▶ `stepi (si)`
Continue to the next instruction
- ▶ `finish`
Execute up to function return
- ▶ `backtrace (bt)`
Display the program stack



GDB advanced commands (1/3)

- ▶ `info threads (i threads)`
Display the list of threads that are available
- ▶ `info breakpoints (i b)`
Display the list of breakpoints/watchpoints
- ▶ `delete <n> (d <n>)`
Delete breakpoint <n>
- ▶ `thread <n> (t <n>)`
Select thread number <n>
- ▶ `frame <n> (f <n>)`
Select a specific frame from the backtrace, the number being the one displayed when using `backtrace` at the beginning of each line



GDB advanced commands (2/3)

- ▶ `watch <variable>` or `watch \*<address>`
Add a watchpoint on a specific variable/address.
- ▶ `print variable = value` (p `variable = value`)
Modify the content of the specified variable with a new value
- ▶ `break foobar.c:42 if condition`
Break only if the specified condition is true
- ▶ `watch <variable> if condition`
Trigger the watchpoint only if the specified condition is true
- ▶ `display <expr>`
Automatically prints expression each time program stops
- ▶ `x/<n><u> <address>`
Display memory at the provided address. `n` is the amount of memory to display, `u` is the type of data to be displayed (b/h/w/g). Instructions can be displayed using the `i` type.



GDB advanced commands (3/3)

- ▶ `list <expr>`
Display the source code associated to the current program counter location.
- ▶ `disassemble <location,start_offset,end_offset> (disas)`
Display the assembly code that is currently executed.
- ▶ `p function(arguments)`
Execute a function using GDB. NOTE: be careful of any side effects that may happen when executing the function
- ▶ `p $newvar = value`
Declare a new gdb variable that can be used locally or in command sequence
- ▶ `define <command_name>`
Define a new command sequence. GDB will prompt for the sequence of commands.



Remote debugging

- ▶ In a non-embedded environment, debugging takes place using `gdb` or one of its front-ends.
- ▶ `gdb` has direct access to the binary and libraries compiled with debugging symbols, which is often false for embedded systems (binaries are stripped, without `debug_info`) to save storage space.
- ▶ For the same reason, embedding the `gdb` program on embedded targets is rarely desirable (2.4 MB on x86).
- ▶ Remote debugging is preferred
 - `ARCH-linux-gdb` is used on the development workstation, offering all its features.
 - `gdbserver` is used on the target system (only 400 KB on arm).


`ARCH-linux-gdb`

`gdbserver`





Remote debugging: architecture





Remote debugging: target setup

- ▶ On the target, run a program through `gdbserver`.
Program execution will not start immediately.
`gdbserver :<port> <executable> <args>`
`gdbserver /dev/ttyS0 <executable> <args>`
- ▶ Otherwise, attach `gdbserver` to an already running program:
`gdbserver --attach :<port> <pid>`
- ▶ You can also start `gdbserver` without passing any program to start or attach (and set the target program later, on client side):
`gdbserver --multi :<port>`



Remote debugging: host setup

- ▶ Then, on the host, start `ARCH-linux-gdb <executable>`, and use the following `gdb` commands:
 - To tell `gdb` where shared libraries are:
`gdb> set sysroot <library-path>` (typically path to build space without `lib/`)
 - To connect to the target:
`gdb> target remote <ip-addr>:<port>` (networking)
`gdb> target remote /dev/ttyUSB0` (serial link)
 - Make sure to replace `target remote` with `target extended-remote` if you have started `gdbserver` with the `--multi` option
 - If you did not set the program to debug on `gdbserver` commandline:
`gdb> set remote exec-file <path_to_program_on_target>`



Coredumps for post mortem analysis

- ▶ It is sometime not possible to have a debugger attached when a crash occurs
- ▶ Fortunately, Linux can generate a `core` file (a snapshot of the whole process memory at the moment of the crash), in the ELF format. `gdb` can use this `core` file to let us analyze the state of the crashed application
- ▶ On the target
 - Use `ulimit -c unlimited` in the shell starting the application, to enable the generation of a `core` file when a crash occurs
 - The output name and path for the coredump file can be modified using `/proc/sys/kernel/core_pattern` (see [man 5 core](#))
 - Example: `echo /tmp/mycore > /proc/sys/kernel/core_pattern`
 - Depending on the system configuration, the `core_pattern` file may be rewritten automatically by some software to handle core files or even disable core generation (eg: `systemd`)
- ▶ On the host
 - After the crash, transfer the `core` file from the target to the host, and run `ARCH-linux-gdb -c core-file application-binary`



- ▶ Coredumps can be huge for complex applications
- ▶ minicoredumper is a userspace tool based on the standard core dump feature
 - Based on the possibility to redirect the core dump output to a user space program via a pipe
- ▶ Based on a JSON configuration file, it can:
 - save only the relevant sections (stack, heap, selected ELF sections)
 - compress the output file
 - save additional information from /proc
- ▶ <https://github.com/diamon/minicoredumper>
- ▶ “Efficient and Practical Capturing of Crash Data on Embedded Systems”
 - Presentation by minicoredumper author John Ogness
 - Video: <https://www.youtube.com/watch?v=q2zmrwrgLJGs>
 - Slides: elinux.org/images/8/81/Eoss2023_ogness_minicoredumper.pdf



GDB: going further

- ▶ Tutorial: Debugging Embedded Devices using GDB - Chris Simmonds, 2020
 - Slides: <https://elinux.org/images/0/01/Debugging-with-gdb-csimmonds-elce-2020.pdf>
 - Video: https://www.youtube.com/watch?v=JGhAgd2a_Ck





GDB Python Extension

- ▶ GDB features a [python integration](#), allowing to script some debugging operations
- ▶ When executing python under GDB, a module named *gdb* is available and all the GDB specific classes are accessible under this module
- ▶ Allows to add new types of commands, breakpoint, printers
 - Used by the kernel to create new commands with the python GDB scripts
- ▶ Allows full control and observability over the debugged program using GDB capabilities from Python scripts
 - Controlling execution, adding breakpoints, watchpoints, etc
 - Accessing the process memory, frames, symbols, etc





GDB Python Extension (1/2)

```
class PrintOpenFD(gdb.FinishBreakpoint):
    def __init__(self, file):
        self.file = file
        super(PrintOpenFD, self).__init__()

    def stop (self):
        print ("---> File " + self.file + " opened with fd " + str(self.return_value))
        return False

class PrintOpen(gdb.Breakpoint):
    def stop(self):
        PrintOpenFD(gdb.parse_and_eval("file").string())
        return False

class TraceFDs (gdb.Command):
    def __init__(self):
        super(TraceFDs, self).__init__("tracefds", gdb.COMMAND_USER)

    def invoke(self, arg, from_tty):
        print("Hooking open() with custom breakpoint")
        PrintOpen("open")

TraceFDs()
```



GDB Python Extension (2/2)

- ▶ Python scripts can be loaded using `gdb source` command
 - Or the script can be named `<program>-gdb.py` and will be loaded automatically by GDB

```
(gdb) source trace_fds.py
(gdb) tracefds
Hooking open() with custom breakpoint
Breakpoint 1 at 0x33e0
(gdb) run
Starting program: /usr/bin/touch foo bar
Temporary breakpoint 2 at 0x5555555587da
---> File foo opened with fd 3
Temporary breakpoint 3 at 0x5555555587da
---> File bar opened with fd 0
```




Common debugging issues

- ▶ You will likely encounter some issues while debugging, like poor address->symbols conversion, "optimized out" values or functions, empty backtraces...
- ▶ A quick checklist before starting debugging can spare you some troubles:
 - Make sure your host binary has **debug symbols**: with gcc, ensure `-g` is provided, and use non-stripped version with host gdb
 - Disable **optimizations** on final binary (`-O0`) if possible, or at least use a less intrusive level (`-Og`)
 - Static functions can for example be folded into caller depending on the optimization level, so they would be missing from backtraces
 - Prevent code optimization from reusing frame pointer register: with GCC, make sure `-fno-omit-frame-pointer` option is set
 - Not only true for debugging: any profiling/tracing tool relying on backtraces will benefit from it
- ▶ Your application is probably composed of multiple libraries: you will need to apply those configurations on all used components!




Debugging an application crash

- ▶ Code generation analysis with compiler-explorer
- ▶ Using GDB and its Python support
- ▶ Analyzing and using a coredump



Application Tracing

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





strace



strace

System call tracer - <https://strace.io>

- ▶ Available on all GNU/Linux systems
Can be built by your cross-compiling toolchain generator or by your build system.
- ▶ Allows to see what any of your processes is doing: accessing files, allocating memory... Often sufficient to find simple bugs.
- ▶ Usage:
`strace <command>` (starting a new process)
`strace -f <command>` (follow child processes too)
`strace -p <pid>` (tracing an existing process)
`strace -c <command>` (time statistics per system call)
`strace -e <expr> <command>` (use expression for advanced filtering)

See [the strace manual](#) for details.



Image credits: <https://strace.io/>



strace example output

```
> strace cat Makefile
[...]  
fstat64(3, {st_mode=S_IFREG|0644, st_size=111585, ...}) = 0  
mmap2(NULL, 111585, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb7f69000  
close(3) = 0  
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)  
open("/lib/tls/i686/cmov/libc.so.6", O_RDONLY) = 3  
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\320h\1\0004\0\0\0\344"... , 512) = 512  
fstat64(3, {st_mode=S_IFREG|0755, st_size=1442180, ...}) = 0  
mmap2(NULL, 1451632, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7e06000  
mprotect(0xb7f62000, 4096, PROT_NONE) = 0  
mmap2(0xb7f66000, 9840, PROT_READ|PROT_WRITE,  
      MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xb7f66000  
close(3) = 0  
[...]  
openat(AT_FDCWD, "Makefile", O_RDONLY) = 3  
newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=173, ...}, AT_EMPTY_PATH) = 0  
fadvise64(3, 0, 0, POSIX_FADV_SEQUENTIAL) = 0  
mmap(NULL, 139264, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xf7f290d28000  
read(3, "ifneq ($(KERNELRELEASE),)\nobj-m "... , 131072) = 173  
write(1, "ifneq ($(KERNELRELEASE),)\nobj-m "... , 173ifneq ($(KERNELRELEASE),)
```

Hint: follow the open file descriptors returned by `open()`. This tells you what files are handled by further system calls.



strace -c example output

```
> strace -c cheese
```

% time	seconds	usecs/call	calls	errors	syscall
36.24	0.523807	19	27017		poll
28.63	0.413833	5	75287	115	ioctl
25.83	0.373267	6	63092	57321	recvmsg
3.03	0.043807	8	5527		writev
2.69	0.038865	10	3712		read
2.14	0.030927	3	10807		getpid
0.28	0.003977	1	3341	34	futex
0.21	0.002991	3	1030	269	openat
0.20	0.002889	2	1619	975	stat
0.18	0.002534	4	568		mmap
0.13	0.001851	5	356		mprotect
0.10	0.001512	2	784		close
0.08	0.001171	3	461	315	access
0.07	0.001036	2	538		fstat

...



ltrace



A tool to trace **shared** library calls used by a program and all the signals it receives

- ▶ Very useful complement to `strace`, which shows only system calls.
- ▶ Of course, works even if you don't have the sources
- ▶ Allows to filter library calls with regular expressions, or just by a list of function names.
- ▶ With the `-S` option it shows system calls too!
- ▶ Also offers a summary with its `-c` option.
- ▶ Manual page: <https://linux.die.net/man/1/ltrace>
- ▶ Works better with *glibc*. `ltrace` used to be broken with *uClibc* (now fixed), and is not supported with *Musl* (Buildroot 2022.11 status).

See <https://en.wikipedia.org/wiki/Ltrace> for details



ltrace example output

```
# ltrace ffmpeg -f video4linux2 -video_size 544x288 -input_format mjpeg -i /dev
/video0 -pix_fmt rgb565le -f fbdev /dev/fb0
__libc_start_main([ "ffmpeg", "-f", "video4linux2", "-video_size"... ] <unfinished ...>
setvbuf(0xb6a0ec80, nil, 2, 0) = 0
av_log_set_flags(1, 0, 1, 0) = 1
strchr("f", ':') = nil
strlen("f") = 1
strncmp("f", "L", 1) = 26
strncmp("f", "h", 1) = -2
strncmp("f", "?", 1) = 39
strncmp("f", "help", 1) = -2
strncmp("f", "-help", 1) = 57
strncmp("f", "version", 1) = -16
strncmp("f", "buildconf", 1) = 4
strncmp("f", "formats", 1) = 0
strlen("formats") = 7
strncmp("f", "muxers", 1) = -7
strncmp("f", "demuxers", 1) = 2
strncmp("f", "devices", 1) = 2
strncmp("f", "codecs", 1) = 3
...
```



ltrace summary

Example summary at the end of the ltrace output (-c option)

% time	seconds	usecs/call	calls	function
52.64	5.958660	5958660	1	__libc_start_main
20.64	2.336331	2336331	1	avformat_find_stream_info
14.87	1.682895	421	3995	strcmp
7.17	0.811210	811210	1	avformat_open_input
0.75	0.085290	584	146	av_freep
0.49	0.055150	434	127	strlen
0.29	0.033008	660	50	av_log
0.22	0.025090	464	54	strcmp
0.20	0.022836	22836	1	avformat_close_input
0.16	0.017788	635	28	av_dict_free
0.15	0.016819	646	26	av_dict_get
0.15	0.016753	440	38	strchr
0.13	0.014536	581	25	memset
...				
100.00	11.318773		4762	total



LD_PRELOAD



Shared libraries

- ▶ Shared libraries are provided as `.so` files that are actually ELF files
 - Loaded at startup by `ld.so` (the dynamic loader)
 - Or at runtime using `dlopen()` from your code
- ▶ When starting a program (an ELF file actually), the kernel will parse it and load the interpreter that needs to be invoked.
 - Most of the time `PT_INTERP` program header of the ELF file is set to `ld-linux.so`.
- ▶ At loading time, the dynamic loader `ld.so` will resolve all the symbols that are present in dynamic libraries.
- ▶ Shared libraries are loaded only once by the OS and then mappings are created for each application that uses the library.
 - This allows to reduce the memory used by libraries.



Hooking Library Calls

- ▶ In order to do some more complex library call hooks, one can use the `LD_PRELOAD` environment variable.
- ▶ `LD_PRELOAD` is used to specify a shared library that will be loaded before any other library by the dynamic loader.
- ▶ Allows to intercept all library calls by preloading another library.
 - Overrides libraries symbols that have the same name.
 - Allows to redefine only a few specific symbols.
 - "Real" symbol can still be loaded and used with `dlsym` ([man 3 dlsym](#))
- ▶ Used by some debugging/tracing libraries (*libsegfault*, *libefence*)
- ▶ Works for C and C++.



LD_PRELOAD example 1/2

- ▶ Library snippet that we want to preload using *LD_PRELOAD*:

```
#include <string.h>
#include <unistd.h>

ssize_t read(int fd, void *data, size_t size) {
    memset(data, 0x42, size);
    return size;
}
```

- ▶ Compilation of the library for *LD_PRELOAD* usage:

```
$ gcc -shared -fPIC -o my_lib.so my_lib.c
```

- ▶ Preloading the new library using *LD_PRELOAD*:

```
$ LD_PRELOAD=./my_lib.so ./exe
```



LD_PRELOAD example 2/2

- ▶ Chaining a call to the real symbol to avoid altering the application behavior:

```
#include <stdio.h>
#include <unistd.h>
#include <dlfcn.h>

ssize_t read(int fd, void *data, size_t size)
{
    size_t (*read_func)(int, void *, size_t);
    char *error;

    read_func = dlsym(RTLD_NEXT, "read");
    if (!read_func) {
        fprintf(stderr, "Can not find read symbol: %s\n", dlerror());
        return 0;
    }
    fprintf(stderr, "Trying to read %lu bytes to %p from file descriptor %d\n", size, data, fd);
    return read_func(fd, data, size);
}
```




uprobes and perf



Probes in linux

- ▶ The linux kernel is able to dynamically add some instrumentation (or "**probes**") to almost any code running on a platform, either in userspace, kernel space, or both.
- ▶ This mechanism works by "patching" the code at runtime to insert the probe. When the patched code is executed, the probe records the execution. It can also collect additional data.
- ▶ There are different kinds of probes exposed by the kernel:
 - **uprobes**: hook on almost any userspace instruction and capture local data
 - **uretprobes**: hook on userspace function exit and capture return value
 - **entry fprobe**: hook on kernel function entry
 - **exit fprobe**: hook on kernel function exit
 - **kprobes**: hook on almost any kernel instruction and capture local data
 - **kretprobe**: hook on kernel function exit and capture return value



- ▶ *uprobe* is a probe mechanism offered by the kernel allowing to trace userspace code.
- ▶ Can target any userspace instruction
 - Internally patches the loaded `.text` section with breakpoints that are handled by the kernel trace system
- ▶ Exposed by file `/sys/kernel/tracing/uprobe_events`
- ▶ User is expected to compute the offset of the targeted instruction inside the corresponding VMA (containing the `.text` section) of the targeted process

```
echo 'p /bin/bash:0x4245c0' > /sys/kernel/tracing/uprobe_events
```

- ▶ Uprobes are wrapped by some common tools (e.g: `perf`, `bcc`) for easier usage
- ▶ `trace/uprobetracer`



The perf tool

- ▶ *perf* tool was started as a tool to profile application under Linux using performance counters ([man 1 perf](#)).
- ▶ It became much more than that and now allows to manage tracepoints, kprobes and uprobes.
- ▶ *perf* can profile both user-space and kernel-space execution.
- ▶ *perf* is based on the `perf_event` interface that is exposed by the kernel.
- ▶ Provides a set of operations, each having specific arguments (see *perf* help).
 - `stat`, `record`, `report`, `top`, `annotate`, `ftrace`, `list`, `probe`, etc



Using perf record

- ▶ *perf record* allows to record performance events per-thread, per-process and per-cpu basis.
- ▶ Kernel needs to be configured with `CONFIG_PERF_EVENTS=y`.
- ▶ This is the first command that needs to be run to gather data from program execution and output them into `perf.data`.
- ▶ `perf.data` file can then be analyzed using `perf annotate` and `perf report`.
 - Useful on embedded systems to analyze data on another computer.



Probing userspace functions

- ▶ List functions that can be probed in a specific executable:

```
$ perf probe --source=<source_dir> -x my_app -F
```

- ▶ List lines number that can be probed in a specific executable/function:

```
$ perf probe --source=<source_dir> -x my_app -L my_func
```

- ▶ Create uprobes on user-space library/executable functions:

```
$ perf probe -x /lib/libc.so.6 printf  
$ perf probe -x my_app my_func:3 my_var  
$ perf probe -x my_app my_func%return \x$retval
```

- ▶ Record the execution of these tracepoints:

```
$ perf record -e probe_my_app:my_func_L3 -e probe_libc:printf
```



Analyzing of application interactions

- ▶ Analyze dynamic library calls from an application using *ltrace*.
- ▶ Overriding a library function with `LD_PRELOAD`.
- ▶ Using *strace* to analyze program syscalls.




Memory Issues

© Copyright 2004-2025, Bootlin.

Creative Commons BY-SA 3.0 license.

Corrections, suggestions, contributions and translations are welcome!





Usual Memory Issues

- ▶ Programming (almost) always involves accessing memory
- ▶ If done incorrectly, a large variety of errors can be triggered
 - Segmentation Faults can happen when accessing invalid memory addresses (NULL pointers or use-after-free for instance)
 - Buffer Overflows can happen if accessing a buffer outside its boundaries
 - Memory Leaks when allocating memory and forgetting to free it after usage
- ▶ Fortunately, there are tools to debug these errors



Segmentation Faults

- ▶ Segmentation Faults are generated by the kernel when a program tries to access a memory area that it is not allowed to or to access it in an incorrect way
 - Might be generated by a write on a read only memory zone
 - Can also be triggered when trying to execute memory that is not executable

```
int *ptr = NULL;  
*ptr = 1;
```

- ▶ Execution will yield a Segmentation fault message in the terminal

```
$ ./program  
Segmentation fault
```



Buffer Overflows

- ▶ Buffer Overflows are easily triggered when accessing an array outside of its boundaries (most often past the end)
- ▶ Such access might generate a crash or not depending on the access
 - Writing past the end of a `malloc()`'ed array will most often overwrite the malloc data structure leading to corruption
 - Writing past the end of an array allocated on the stack can corrupt data on the stack
 - Reading past the end of an array might generate a segfault but not always, this depends on the area of memory that is accessed

```
uint32_t *array = malloc(10 * sizeof(*array));  
array[10] = 0xDEADBEEF;
```



Memory Leaks

- ▶ Memory leaks are another class of memory errors that will not directly trigger a crash but will exhaust the system memory (sooner or later)
- ▶ This happens when allocating memory in your program and not releasing it after using it
- ▶ Can trigger in production when the program runs for a very long time
 - Better to debug that kind of problem early in the development process

```
void func1(void) {  
    uint32_t *array = malloc(10 * sizeof(*array));  
    do_something_with_array(array);  
}
```



Valgrind memcheck



Valgrind (1/2)

- ▶ **Valgrind** is an instrumentation framework for building dynamic analysis tools
- ▶ `valgrind` is also a tool that is based on this framework and provides a memory error detector, heap profilers and others profilers.
- ▶ It is supported on all the popular platforms: Linux on x86, x86_64, arm (armv7 only), arm64, mips32, s390, ppc32 and ppc64.





Valgrind (2/2)

- ▶ Works by adding its own instrumentation to your code and then running in on its own virtual cpu core. Significantly slows down execution, and thus is suited for debugging and profiling
- ▶ *Memcheck* is the default *valgrind* tool and it detects memory-management errors
 - Access to invalid memory zones, use of uninitialized values, memory leaks, bad freeing of heap blocks, etc
 - Can be run on any application, no need to recompile them



```
$ valgrind --tool=memcheck --leak-check=full <program>
```



Valgrind Memcheck usage and report

```
$ valgrind ./mem_leak
==202104== Memcheck, a memory error detector
==202104== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==202104== Using Valgrind-3.18.1 and LibVEX; rerun with -h for copyright info
==202104== Command: ./mem_leak
==202104==
==202104== Conditional jump or move depends on uninitialised value(s)
==202104==    at 0x109161: do_actual_jump (in /home/user/mem_leak)
==202104==    by 0x109187: compute_address (in /home/user/mem_leak)
==202104==    by 0x1091A2: do_jump (in /home/user/mem_leak)
==202104==    by 0x1091D7: main (in /home/user/mem_leak)
==202104==
==202104== HEAP SUMMARY:
==202104==    in use at exit: 120 bytes in 1 blocks
==202104==    total heap usage: 1 allocs, 0 frees, 120 bytes allocated
==202104==
==202104== LEAK SUMMARY:
==202104==    definitely lost: 120 bytes in 1 blocks
==202104==    indirectly lost: 0 bytes in 0 blocks
==202104==    possibly lost: 0 bytes in 0 blocks
==202104==    still reachable: 0 bytes in 0 blocks
==202104==    suppressed: 0 bytes in 0 blocks
==202104== Rerun with --leak-check=full to see details of leaked memory
```




Valgrind and VGDB

- ▶ Valgrind can also act as a GDB server which can receive and process commands. One can interact with valgrind gdb server either with a gdb client, or directly with `vgdb` program (provided with valgrind). `vgdb` can be used in different ways:
 - As a standalone CLI program to send "monitor" commands to valgrind
 - As a relay between a gdb client and an existing valgrind session
 - As a server to drive multiple valgrind sessions from a remote gdb client
- ▶ See `man 1 vgdb` for available modes, commands and options



Using GDB with Memcheck

- ▶ *valgrind* allows to attach with GDB to the process that is currently analyzed.

```
$ valgrind --tool=memcheck --leak-check=full --vgdb=yes --vgdb-error=0 ./mem_leak
```

- ▶ Then attach gdb to the valgrind gdbserver using *vgdb*

```
$ gdb ./mem_leak  
(gdb) target remote | vgdb
```

- ▶ If valgrind detects an error, it will stop the execution and break into GDB.

```
(gdb) continue  
Continuing.  
  
Program received signal SIGTRAP, Trace/breakpoint trap.  
0x000000000109161 in do_actual_jump (p=0x4a52040) at mem_leak.c:5  
5   if (p[1])  
(gdb) bt  
#0  0x000000000109161 in do_actual_jump (p=0x4a52040) at mem_leak.c:5  
#1  0x000000000109188 in compute_address (p=0x4a52040) at mem_leak.c:11  
#2  0x0000000001091a3 in do_jump (p=0x4a52040) at mem_leak.c:16  
#3  0x0000000001091d8 in main () at mem_leak.c:27
```



Electric Fence



libefence (1/2)

- ▶ *libefence* is more lightweight than *valgrind* but less precise
- ▶ Allows to catch two types of common memory errors
 - Buffer overflows and use after free
- ▶ *libefence* will actually trigger a segfault upon the first error encountered in order to generate a coredump.
- ▶ Uses a shared library that can either be linked with statically (`-lefence`) or preloaded using `LD_PRELOAD`.

```
$ gcc -g program.c -o program
$ LD_PRELOAD=libefence.so.0.0 ./program
```

```
Electric Fence 2.2 Copyright (C) 1987-1999 Bruce Perens <bruce@perens.com>
Segmentation fault (core dumped)
```



libefence (2/2)

- ▶ Upon segfault, a coredump will be generated in the current directory
- ▶ This coredump can be opened with GDB and will pinpoint the exact location where the error happened

```
$ gdb ./program core-program-3485
Reading symbols from ./libefence...
[New LWP 57462]
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Core was generated by `./libefence'.
Program terminated with signal SIGSEGV, Segmentation fault.
#0  main () at libefence.c:8
      data[99] = 1;
(gdb)
```



Practical lab - Debugging Memory Issues




Debug various memory issues using specific tooling

- ▶ Memory leak and misbehavior detection with *valgrind* and *vgdb*.



Application Profiling

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





- ▶ Profiling is the act of gathering data from a program execution in order to analyze them and then optimize or fix performance issues.
- ▶ Profiling is achieved by using programs that insert instrumentation in the code or leverage kernel/userspace mechanisms.
 - Profiling function calls and count of calls allow to optimize performance.
 - Profiling processor usage allows to optimize performance and reduce power usage.
 - Profiling memory usage allows to optimize memory consumption.
- ▶ After profiling, the data set must be analyzed to identify potential improvements (and not the reverse!).



"Premature optimization is the root of all evil", Donald Knuth

- ▶ Profiling is often useful to identify and fix performance issues.
- ▶ Performance can be affected by memory usage, IOs overload, or CPU usage.
- ▶ Gathering profiling data before trying to fix performance issues is needed to do the correct choices.
- ▶ Profiling is often guided by a first coarse-grained analysis using some classic tools.
- ▶ Once the class of problems has been identified, a fine-grained profiling analysis can be done.



Profiling metrics



- ▶ Multiple tools allow to profile various metrics.
- ▶ Memory usage with *Massif*, *heaptrack* or *memusage*.
- ▶ Function calls using *perf* and *callgrind*.
- ▶ CPU hardware usage (Cache, MMU, etc) using *perf*.
- ▶ Profiling data can include both the user space application and kernel.



Memory profiling



Memory profiling

- ▶ Profiling memory usage (heap/stack) in an application is useful for optimization.
- ▶ Allocating too much memory can lead to system memory exhaustion.
- ▶ Allocating/freeing memory too often can lead to the kernel spending a considerable amount of time in `clear_page()`.
 - The kernel clears pages before giving them to processes to avoid data leakage.
- ▶ Reducing application memory footprint can allow optimizing cache usage as well as page miss.



Massif usage

- ▶ *Massif* is a tool provided by *valgrind* which allows to profile heap usage during the program execution (user-space only).
- ▶ Works by making snapshots of allocations.

```
$ valgrind --tool=massif --time-unit=B program
```


- ▶ Once executed, a *massif.out.<pid>* file will be generated in the current directory
- ▶ *ms_print* tool can then be used to display a graph of heap allocation

```
$ ms_print massif.out.275099
```

- ▶ #: Peak allocation
- ▶ @: Detailed snapshot (count can be adjusted thanks to *--detailed-freq*)



Massif report




Number of snapshots: 52

Detailed snapshots: [9, 19, 22 (peak), 32, 42]



massif-visualizer - Visualizing massif profiling data





heaptrack usage


- ▶ *heaptrack* is a heap memory profiler for Linux.
 - Works with LD_PRELOAD library.
- ▶ Finer tracking than with Massif and visualizing tool is more advanced.
 - Each allocation is associated to a stacktrace.
 - Allows finding memory leaks, allocation hotspots and temporary allocations.
- ▶ Results can be seen using GUI (*heaptrack_gui*) or CLI tool (*heaptrack_print*).
- ▶ <https://github.com/KDE/heaptrack>

```
$ heaptrack program
```

- ▶ This will generate a *heaptrack.<process_name>.<pid>.zst* file that can be analyzed using *heaptrack_gui* on another computer.



heaptrack_gui - Visualizing heaptrack profiling data






heaptrack_gui - Flamegraph view





- ▶ memusage is a program that leverages `libmemusage.so` to profile memory usage (`man 1 memusage`) (user-space only).
- ▶ Can profile heap, stack and also mmap memory usage.
- ▶ Profiling information can be shown on the console, logged to a file for post-treatment or visualized in a PNG file.
- ▶ Lightweight solution compared to valgrind *Massif* tool since it uses the `LD_PRELOAD` mechanism.





memusage usage

```
$ memusage convert foo.png foo.jpg
Memory usage summary: heap total: 2635857, heap peak: 2250856, stack peak: 83696
      total calls  total memory  failed calls
malloc|      1496      2623648          0
realloc|         6        3744          0 (nomove:0, dec:0, free:0)
calloc|         16       8465          0
  free|      1480      2521334
Histogram for block sizes:
  0-15          329  21% =====
 16-31          239  15% =====
 32-47          287  18% =====
 48-63          321  21% =====
 64-79           43   2% =====
 80-95          141   9% =====
...
21424-21439         1 <1%
32768-32783         1 <1%
32816-32831         1 <1%
  large           3 <1%
```



Execution profiling



- ▶ In order to optimize a program, one may have to understand what hardware resources are used.
- ▶ Many hardware elements can have an impact on the program execution:
 - CPU cache performance can be degraded by an application without memory spatial locality.
 - Page miss due to using too much memory without spatial locality.
 - Alignment faults when doing misaligned accesses.



Using *perf stat*

- ▶ `perf stat` allows to profile an application by gathering performance counters.
 - Using performance counters might require *root* permissions. This can be modified using `# echo -1 > /proc/sys/kernel/perf_event_paranoid`
- ▶ The number of performance counters that are present on the hardware are often limited.
- ▶ Requesting more events than possible will result in multiplexing and `perf` will scale the results.
- ▶ Collected performance counters are then approximate.
 - To acquire more precise numbers, reduce the number of events observed and run `perf` multiple times changing the events set to observe all the expected events.
 - See [perf wiki](#) for more information.



perf stat example (1/2)

```
$ perf stat convert foo.png foo.jpg
```

```
Performance counter stats for 'convert foo.png foo.jpg':
```

45,52 msec	task-clock	#	1,333 CPUs utilized	
4	context-switches	#	87,874 /sec	
0	cpu-migrations	#	0,000 /sec	
1 672	page-faults	#	36,731 K/sec	
146 154 800	cycles	#	3,211 GHz	(81,16%)
6 984 741	stalled-cycles-frontend	#	4,78% frontend cycles idle	(91,21%)
81 002 469	stalled-cycles-backend	#	55,42% backend cycles idle	(91,36%)
222 687 505	instructions	#	1,52 insn per cycle	
		#	0,36 stalled cycles per insn	(91,21%)
37 776 174	branches	#	829,884 M/sec	(74,51%)
567 408	branch-misses	#	1,50% of all branches	(70,62%)

```
0,034156819 seconds time elapsed
```

```
0,041509000 seconds user
```

```
0,004612000 seconds sys
```

- *NOTE: the percentage displayed at the end denotes the time during which the kernel measured the event due to multiplexing*



perf stat example (2/2)

► List all events:

```
$ perf list
List of pre-defined events (to be used in -e):

branch-instructions OR branches      [Hardware event]
branch-misses                        [Hardware event]
cache-misses                         [Hardware event]
cache-references                     [Hardware event]
...
```

► Count *L1-dcache-load-misses* and *branch-load-misses* events for a specific command

```
$ perf stat -e L1-dcache-load-misses,branch-load-misses cat /etc/fstab
...
Performance counter stats for 'cat /etc/fstab':

23 418      L1-dcache-load-misses
 7 192      branch-load-misses
...
```



Cachegrind


- ▶ *Cachegrind* is a tool provided by *valgrind* for profiling program interactions with the instruction and data cache hierarchy.
 - *Cachegrind* also profiles branch prediction success.
- ▶ Simulate a machine with independent I\$ and D\$ backed with a unified L2 cache.
- ▶ Really helpful to detect cache usage problems (too many misses, etc).

```
$ valgrind --tool=cachegrind --cache-sim=yes ./my_program
```

- ▶ It generates a `cachegrind.out.<pid>` file containing the measures
- ▶ `cg_annotate` is a CLI tool used to visualize cachegrind simulation results.
- ▶ It also has a `--diff` option to allow comparing two measures files



Kcachegrind - Visualizing Cachegrind profiling data





Callgrind


- ▶ Provided by *valgrind* and allowing to profile an application call graph (user-space only).
- ▶ Collects the number of instructions executed during your program execution and associate these data with the source lines
- ▶ Records the call relationship between functions and their call count.

```
$ valgrind --tool=callgrind ./my_program
```

- ▶ `callgrind_annotate` is a CLI tool used to visualize callgrind simulation results.
- ▶ `Kcachegrind` can visualize *callgrind* results too.
- ▶ The cache simulation (done using `cachegrind`) has some accuracy shortcomings (See [Cachegrind accuracy](#))



Kcachegrind - Visualizing Callgrind profiling data






Profiling an application using various tools

- ▶ Profiling application heap using *Massif*.
- ▶ Profiling an application with *Cachegrind*, *Callgrind* and *KCachegrind*.
- ▶ Analyzing application performance with *perf*.



System-wide Profiling & Tracing

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





System-wide Profiling & Tracing

- ▶ Sometimes, the problems are not tied to an application but rather due to the usage of multiple layers (drivers, application, kernel).
- ▶ In that case, it might be useful to analyze the whole stack.
- ▶ The kernel already includes a large number of tracepoints that can be recorded using specific tools.
- ▶ New tracepoints can also be created statically or dynamically using various mechanisms (kprobes for instance).



kprobes



Kprobes

- ▶ Allows to insert breaks at almost any kernel address dynamically and to extract debugging and performance information
- ▶ Uses code patching to modify text code to insert calls to specific handlers
 - `kprobes` allows to execute specific handlers when the hooked instruction is executed
 - `kretprobes` will trigger when returning from a function allowing to extract the return value of functions but also display the parameters that were used for the function call
- ▶ Needs some basic kernel configuration:
 - `CONFIG_KPROBES=y` to enable general kprobe support
 - `CONFIG_KALLSYMS_ALL=y` to allow hooking probes using `<symbol_name>` instead of raw function address
 - `CONFIG_KPROBE_EVENTS=y` to enable kprobes usage as tracing events in `tracefs`
- ▶ At the lowest level, `k(ret)probes` are manipulated with dedicated kernel APIs, allowing to write our own kprobe tools (eg as kernel modules)
- ▶ Can also be used from userspace with `/sys/kernel/tracing/kprobe_events`
- ▶ See [trace/kprobes](#) for more information



Basic kprobe tracing (1/2)

- ▶ Add a kprobe on `do_sys_openat2`:

```
$ echo "p:my_probe do_sys_openat2" > /sys/kernel/tracing/kprobe_events
```

- ▶ Add a kprobe in the same function but at a specific offset, and capture some arguments

```
$ echo "p:my_probe_2 do_sys_openat2+0x7c file=%r2" > /sys/kernel/tracing/kprobe_events
```

- ▶ Insert a kretprobe

```
$ echo 'r:my_retprobe do_sys_openat2 $retval' > /sys/kernel/tracing/kprobe_events
```



Basic kprobe tracing (2/2)

- ▶ Show existing kprobes

```
$ cat /sys/kernel/tracing/kprobe_events
```

- ▶ Enable a kprobe (ie: start capturing the corresponding event)

```
$ echo 1 > /sys/kernel/tracing/events/kprobes/my_probe/enable
```

- ▶ Get data emitted by kprobes

```
$ cat /sys/kernel/tracing/trace
```

- ▶ Delete a kprobe

```
$ echo "-:my_probe" >> /sys/kernel/tracing/kprobe_events
```



System-wide Profiling & Tracing

perf

- ▶ *perf* allows to do a wide range of tracing and recording operations.
- ▶ The kernel already contains events and tracepoints that can be used. The list is given using `perf list`.
- ▶ Syscall tracepoints should be enabled in kernel configuration using `CONFIG_FTRACE_SYSCALLS`.
- ▶ New tracepoint can be created dynamically on all symbols and registers when debug info are not present.
- ▶ Tracing functions, recording variables and parameters content using their names will require a kernel compiled with `CONFIG_DEBUG_INFO`.
- ▶ If *perf* does not find `vmlinux` you have to provide it using `-k <vmlinux>`.



perf example

- ▶ List all events that matches `syscalls:*`

```
$ perf list syscalls:*
```

```
List of pre-defined events (to be used in -e):
```

<code>syscalls:sys_enter_accept</code>	<code>[Tracepoint event]</code>
<code>syscalls:sys_enter_accept4</code>	<code>[Tracepoint event]</code>
<code>syscalls:sys_enter_access</code>	<code>[Tracepoint event]</code>
<code>syscalls:sys_enter_adjtimex_time32</code>	<code>[Tracepoint event]</code>
<code>syscalls:sys_enter_bind</code>	<code>[Tracepoint event]</code>

```
...
```

- ▶ Record all `syscalls:sys_enter_read` events for `sha256sum` command into `perf.data` file.

```
$ perf record -e syscalls:sys_enter_read sha256sum /bin/busybox
```

```
[ perf record: Woken up 1 times to write data ]
```

```
[ perf record: Captured and wrote 0.018 MB perf.data (215 samples) ]
```



perf report example

- Display the collected samples ordered by time spent.

```
$ perf report
Samples: 591 of event 'cycles', Event count (approx.): 393877062
Overhead Command Shared Object Symbol
 22,88% firefox-esr [nvidia] [k] _nv031568rm
  3,21% firefox-esr ld-linux-x86-64.so.2 [.] __minimal_realloc
  2,00% firefox-esr libc.so.6 [.] __stpncpy_sse3
  1,86% firefox-esr libglib-2.0.so.0.7400.0 [.] g_hash_table_lookup
  1,62% firefox-esr ld-linux-x86-64.so.2 [.] _dl_strtoul
  1,56% firefox-esr [kernel.kallsyms] [k] clear_page_rep
  1,52% firefox-esr libc.so.6 [.] __strncpy_sse2_unaligned
  1,37% firefox-esr ld-linux-x86-64.so.2 [.] strncmp
  1,30% firefox-esr firefox-esr [.] malloc
  1,27% firefox-esr libc.so.6 [.] __GI___strcasecmp_l_sse3
  1,23% firefox-esr [nvidia] [k] _nv013165rm
  1,09% firefox-esr [nvidia] [k] _nv007298rm
  1,03% firefox-esr [kernel.kallsyms] [k] unmap_page_range
  0,91% firefox-esr ld-linux-x86-64.so.2 [.] __minimal_free
```




- ▶ *perf* allows to create dynamic tracepoints on both kernel functions and user-space functions.
- ▶ In order to be able to insert probes, `CONFIG_KPROBES` must be enabled in the kernel.
 - Note: *libelf* is required to compile *perf* with *probe* command support.
- ▶ New dynamic probes can be created and then used using *perf record*.
- ▶ Often on embedded platforms, `vmlinux` is not present on the target and thus only symbols and registers can be used.



perf probe examples (1/3)

- ▶ List all the kernel symbols that can be probed (no debug info needed):

```
$ perf probe --funcs
```

- ▶ Create a new probe on `do_sys_openat2` with *filename* named parameter (debug info required).

```
$ perf probe --vmlinux=vmlinux_file do_sys_openat2 filename:string
```

```
Added new event:
```

```
probe:do_sys_openat2 (on do_sys_openat2 with filename:string)
```

- ▶ Execute `tail` and capture previously created probe event:

```
$ perf record -e probe:do_sys_openat2 tail /var/log/messages
```

```
...
```

```
[ perf record: Woken up 1 times to write data ]
```

```
[ perf record: Captured and wrote 0.003 MB perf.data (19 samples) ]
```



perf probe examples (2/3)

- ▶ Display the recorded tracepoints with *perf script*:

```
$ perf script
tail 164 [000] 3552.956573: probe:do_sys_openat2: (c02c3750) filename_string="/etc/ld.so.cache"
tail 164 [000] 3552.956642: probe:do_sys_openat2: (c02c3750) filename_string="/lib/tls/v7l/neon/vfp/libresolv.so.2"
...
```

- ▶ Create a new probe to capture the return value from `ksys_read`

```
$ perf probe ksys_read%return \${retval}
```

- ▶ Execute `sha256sum` and capture previously created probe events:

```
$ perf record -e probe:ksys_read__return sha256sum /etc/fstab
```



perf probe examples (3/3)

- ▶ List all probes that have been created:

```
$ perf probe -l  
probe:ksys_read__return (on ksys_read%return with ret)
```

- ▶ Remove an existing tracepoint:

```
$ perf probe -d probe:ksys_read__return
```



perf record example

- ▶ Record all events for all cpus (system-wide mode):

```
$ perf record -a  
^C
```

- ▶ Display recorded events from perf.data using perf script

```
$ perf script  
...  
klogd 85 [000] 208.609712: 116584 cycles: b6dd551c memset+0x2c (/lib/libc.so.6)  
klogd 85 [000] 208.609898: 121267 cycles: c0a44c84 _raw_spin_unlock_irq+0x34 (vmlinux)  
klogd 85 [000] 208.610094: 127434 cycles: c02f3ef4 kmem_cache_alloc+0xd0 (vmlinux)  
perf 130 [000] 208.610311: 132915 cycles: c0a44c84 _raw_spin_unlock_irq+0x34 (vmlinux)  
perf 130 [000] 208.619831: 143834 cycles: c0a44cf4 _raw_spin_unlock_irqrestore+0x3c (vmlinux)  
klogd 85 [000] 208.620048: 143834 cycles: c01a07f8 syslog_print+0x170 (vmlinux)  
klogd 85 [000] 208.620241: 126328 cycles: c0100184 vector_swi+0x44 (vmlinux)  
klogd 85 [000] 208.620434: 128451 cycles: c096f228 unix_dgram_sendmsg+0x46c (vmlinux)  
kworker/0:2-mm_ 44 [000] 208.620653: 133104 cycles: c0a44c84 _raw_spin_unlock_irq+0x34 (vmlinux)  
perf 130 [000] 208.620859: 138065 cycles: c0198460 lock_acquire+0x184 (vmlinux)  
...
```



Using perf trace

- ▶ perf trace captures and displays all tracepoints/events that have been triggered when executing a command

```
$ perf trace -e "net:*" ping -c 1 192.168.1.1
PING 192.168.1.1 (192.168.1.1) 56(84) bytes of data.
 0.000 ping/37820 net:net_dev_queue(skbaddr: 0xffff97bbc6a17900, len: 98,
    name: "enp34s0")
 0.005 ping/37820 net:net_dev_start_xmit(name: "enp34s0",
    skbaddr: 0xffff97bbc6a17900, protocol: 2048, len: 98,
    network_offset: 14, transport_offset_valid: 1, transport_offset: 34)
 0.009 ping/37820 net:net_dev_xmit(skbaddr: 0xffff97bbc6a17900, len: 98,
    name: "enp34s0")
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=0.867 ms
```



Using perf top

- ▶ perf top allows to do a live analysis of the running kernel
- ▶ It will sample all function calls and display them ordered by most time consuming one.
- ▶ This allows to profile the whole system usage

```
$ perf top
Samples: 19K of event 'cycles', 4000 Hz, Event count (approx.): 4571734204 lost: 0/0 drop: 0/0
Overhead Shared Object          Symbol
 2,01% [nvidia]                  [k] _nv023368rm
 0,94% [kernel]                  [k] __static_call_text_end
 0,89% [vdso]                    [.] 0x0000000000000655
 0,81% [nvidia]                  [k] _nv027733rm
 0,79% [kernel]                  [k] clear_page_rep
 0,76% [kernel]                  [k] psi_group_change
 0,70% [kernel]                  [k] check_preemption_disabled
 0,69% code                      [.] 0x000000000623108f
 0,60% code                      [.] 0x0000000006231083
 0,59% [kernel]                  [k] preempt_count_add
 0,54% [kernel]                  [k] module_get_kallsym
 0,53% [kernel]                  [k] copy_user_generic_string
```



Using a GUI to display perf data

- ▶ `perf report` is the default way to display perf data, directly in the console
- ▶ There are also graphical tools to display perf data:
 - **Flamegraphs**
 - Visualization based on hierarchical stacks
 - Allows to quickly find bottlenecks and explore the call stack
 - Popularized by Brendan Gregg tools which allows to generate flamegraphs from perf results.
 - **Hotspot** software
 - Developed and maintained by KDAB
 - A larger tool able to generate various types of visualizations from a `perf.data` file
 - Can also perform the actual perf recording



Visualizing data with flamegraphs

- ▶ Get the flamegraph scripts:

```
git clone https://github.com/brendangregg/FlameGraph fl
```

- ▶ Capture data:

```
perf record -g -- sleep 30
```

- The `-g` option records call stacks for each sample

- ▶ Format the data:

```
perf script | ./fl/stackcollapse-perf.pl > out.perf-folded
```

- Other data sources are supported (eg: DTrace, SystemTap, Intel VTune, gdb...)


- ▶ Generate the Flamegraph:

```
./fl/flamegraph.pl out.perf-folded > flamegraph.svg
```

- ▶ The flamegraph can then be opened in a web browser



Flamegraph example: CPU flamegraph



- ▶ The plates on top represent the functions sampled by perf during the recording
- ▶ The plates width represents how often a function has been sampled by perf
- ▶ The plates below represent the call stacks for the sampled functions
- ▶ Flamegraphs are interactive: clicking on a plate will zoom on the corresponding callstack
- ▶ Colors can be tuned at flamegraph generation (eg: to get a clear split between kernel and userspace)




Visualizing data with hotspot (1/2)

- ▶ Designed to provide a frontend to perf data files
- ▶ Can generate flamegraphs on the fly, but not only:
 - CPU/tasks timelines
 - Interactive callstacks navigation
 - Code disassembly
- ▶ Configurable (eg: allows to set paths to find all needed debug informations)



Visualizing data with hotspot (2/2)





ftrace and trace-cmd



- ▶ *ftrace* is a tracing framework within the kernel which stands for "Function Tracer".
- ▶ It offers a wide range of tracing capabilities allowing to observe the system behavior.
 - Trace static tracepoints already inserted at various locations in the kernel (scheduler, interrupts, etc).
 - Relies on GCC `mcount()` capability and kernel code patching mechanism to call *ftrace* tracing handlers.
- ▶ All traces are recorded in a ring buffer that is optimized for tracing.
- ▶ Uses *tracefs* filesystem to control and display tracing events.
 - `# mount -t tracefs nodev /sys/kernel/tracing.`
- ▶ *ftrace* support must be enabled in the kernel using `CONFIG_FTRACE=y`.
- ▶ `CONFIG_DYNAMIC_FTRACE` allows to have a zero overhead tracing support.



- ▶ *ftrace* controls are exposed through some specific files located under `/sys/kernel/tracing`.
 - `current_tracer`: Current tracer that is used.
 - `available_tracers`: List of available tracers that are compiled in the kernel.
 - `tracing_on`: Enable/disable tracing.
 - `trace`: Acquired trace in human readable format. Format will differ depending on the tracer used.
 - `trace_pipe`: same as `trace`, but each read consumes the trace as it is read.
 - `trace_marker{[_raw]}`: Emit comments from userspace in the trace buffer.
 - `set_ftrace_filter`: Filter some specific functions.
 - `set_graph_function`: Graph only the specified functions child.
- ▶ Many other files are exposed, see [trace/ftrace](#).
- ▶ *trace-cmd* CLI and *Kernelshark* GUI tools allow to record and visualize tracing data more easily.



- ▶ ftrace provides several "tracers" which allow to trace different things.
- ▶ The tracer to be used should be written to the `current_tracer` file
 - `nop`: Trace nothing, used to disable all tracing.
 - `function`: Trace all kernel functions that are called.
 - `function_graph`: Similar to `function` but traces both entry and exit.
 - `hwlat`: Trace hardware latency.
 - `irqsoff`: Trace sections where interrupts are disabled.
 - `branch`: Trace likely()/unlikely() prediction errors.
 - `mmiotrace`: Trace all accesses to the hardware (`read[bwlq]/write[bwlq]`).
- ▶ **Warning: Some tracers can be expensive!**

```
# echo "function" > /sys/kernel/tracing/current_tracer
```




function_graph tracer report example

- ▶ The *function_graph* traces all the function that executed and their associated callgraphs
- ▶ Will display the process, CPU, timestamp and function graph:


```
$ trace-cmd report
```

```
...
dd-113  [000]  304.526590: funcgraph_entry:      | sys_write() {
dd-113  [000]  304.526597: funcgraph_entry:      |   ksys_write() {
dd-113  [000]  304.526603: funcgraph_entry:      |     __fdget_pos() {
dd-113  [000]  304.526609: funcgraph_entry:      |       __fget_light();
dd-113  [000]  304.526621: funcgraph_exit:      + 18.500 us |   }
dd-113  [000]  304.526627: funcgraph_entry:      |   vfs_write() {
dd-113  [000]  304.526634: funcgraph_entry:      |     rw_verify_area();
dd-113  [000]  304.526646: funcgraph_entry:      |     write_null();
dd-113  [000]  304.526658: funcgraph_entry:      |     __fsnotify_parent();
dd-113  [000]  304.526669: funcgraph_exit:      + 43.042 us |   }
dd-113  [000]  304.526675: funcgraph_exit:      + 78.833 us | }
dd-113  [000]  304.526680: funcgraph_exit:      + 91.291 us | }
dd-113  [000]  304.526689: funcgraph_entry:      | sys_read() {
dd-113  [000]  304.526695: funcgraph_entry:      |   ksys_read() {
dd-113  [000]  304.526702: funcgraph_entry:      |     __fdget_pos() {
dd-113  [000]  304.526708: funcgraph_entry:      |       __fget_light();
dd-113  [000]  304.526719: funcgraph_exit:      + 18.083 us |   }
```



irqsoff tracer

- ▶ `ftrace irqsoff` tracer allows to trace the irq latency due to interrupts being disabled for too long.
- ▶ Helpful to find why interrupts have high latencies on a system.
- ▶ This tracer will record the longest trace with interrupts being disabled.
- ▶ This tracer needs to be enabled with `CONFIG_IRQSOFF_TRACER=y`.
 - `preemptoff`, `preemptirqsoff` tracers also exist to trace section of code where preemption is disabled.






irqsoff: report example

```
# latency: 276 us, #104/104, CPU#0 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:2)
# -----
# | task: stress-ng-114 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: __irq_usr
# => ended at:  irq_exit
#
#
#           _-----> CPU#
#           / _-----> irqsoff
#           | / _-----> need-resched
#           || / _-----> hardirq/softirq
#           ||| / _-----> preempt-depth
#           |||| /      delay
# cmd      pid  |||| time | caller
#  \      /  |||| \   | /
stress-n-114 0d... 2us : __irq_usr
stress-n-114 0d... 7us : gic_handle_irq <-__irq_usr
stress-n-114 0d... 10us : __handle_domain_irq <-gic_handle_irq
...
stress-n-114 0d... 270us : __local_bh_disable_ip <-__do_softirq
stress-n-114 0d.s. 275us : __do_softirq <-irq_exit
stress-n-114 0d.s. 279us+: tracer_hardirqs_on <-irq_exit
stress-n-114 0d.s. 290us : <stack trace>
```



Hardware latency detector

- ▶ `ftrace hwlat` tracer will help to find if the hardware generates latency.
 - System Management interrupts for instance are non maskable and directly trigger some firmware support feature, suspending CPU execution.
 - Interrupts handled by secure monitor can also cause this kind of latency.
- ▶ If some latency is found with this tracer, the system is probably not suitable for real time usage.
- ▶ Uses a single core looping while interrupts are disabled and measuring the time elapsed between two consecutive time reads.
- ▶ Needs to be builtin the kernel with `CONFIG_HWLAT_TRACER=y`.





- ▶ *trace-cmd* is a tool written by Steven Rostedt which allows interacting with *ftrace* ([man 1 trace-cmd](#)).
- ▶ The tracers supported by *trace-cmd* are those exposed by *ftrace*.
- ▶ *trace-cmd* offers multiple commands:
 - *list*: List available plugins/events that can be recorded.
 - *record*: Record a trace into the file `trace.dat`.
 - *report*: Display `trace.dat` acquisition results.



trace-cmd examples (1/3)

▶ List available tracers

```
$ trace-cmd list -t
blk mmiotrace function_graph function nop
```

▶ List available events

```
$ trace-cmd list -e
...
migrate:mm_migrate_pages_start
migrate:mm_migrate_pages
tlb:tlb_flush
syscalls:sys_exit_process_vm_writev
...
```

▶ List available functions for filtering with function and function_graph tracers

```
$ trace-cmd list -f
...
wait_for_initramfs
__ftrace_invalid_address___64
calibration_delay_done
calibrate_delay
...
```



trace-cmd examples (2/3)

- ▶ Start the function tracer and record data globally on the system

```
$ trace-cmd record -p function
```

- ▶ Use the function graph tracer but filter only `spi_*` functions

```
$ trace-cmd record -l spi_* -p function_graph
```

- ▶ Run the *irqsoff* tracer on the system:

```
$ trace-cmd record -p irqsoff
```

- ▶ Record only `irq_handler_exit/irq_handler_entry` events on the system:

```
$ trace-cmd record -e irq:irq_handler_exit -e irq:irq_handler_entry
```



trace-cmd examples (3/3)

- ▶ Visualize the data that have been acquired in `trace.dat`:

```
$ trace-cmd report
```

- ▶ Reset all the *ftrace* buffers and remove tracers

```
$ trace-cmd reset
```




Remote tracing with trace-cmd

- ▶ *trace-cmd* output can be quite big and thus difficult to store on an embedded platform with limited storage.
- ▶ For that purpose, a `listen` command is available and allows sending the acquisitions over the network:
 - Run `trace-cmd listen -p 6578` on the remote system that will be collecting the traces
 - On the target system, use `trace-cmd record -N <target_ip>:6578` to specify the remote system that will collect the traces





trace_printk()

- ▶ `trace_printk()` allows to emit strings in the trace buffer
- ▶ Useful to trace some specific conditions in your code and display it in the trace buffer

```
#include <linux/ftrace.h>
void read_hw()
{
    if (condition)
        trace_printk("Condition is true!\n");
}
```

- ▶ Will display the following in the trace buffer for function_graph tracer

```
1)          |          read_hw() {
1)          |          /* Condition is true! */
1)  2.657 us |          }
```



Adding ftrace tracepoints (1/2)

- ▶ For some custom needs, it might be needed to add custom tracepoints
- ▶ First, one needs to declare the tracepoint definition in a .h file

```
#undef TRACE_SYSTEM
#define TRACE_SYSTEM subsys

#if !defined(_TRACE_SUBSYS_H) || defined(TRACE_HEADER_MULTI_READ)
#define _TRACE_SUBSYS_H

#include <linux/tracepoint.h>

DECLARE_TRACE(subsys_eventname,
               TP_PROTO(int firstarg, struct task_struct *p),
               TP_ARGS(firstarg, p));

#endif /* _TRACE_SUBSYS_H */

/* This part must be outside protection */
#include <trace/define_trace.h>
```



Adding ftrace tracepoints (2/2)

- ▶ Then, emit tracepoint in a .c file using that header file

```
#include <trace/events/subsys.h>

#define CREATE_TRACE_POINTS
DEFINE_TRACE(subsys_eventname);


void any_func(void)
{
    ...
    trace_subsys_eventname(arg, task);
    ...
}
```

- ▶ See [trace/tracepoints](#) for more information



- ▶ Kernelshark is a Qt-based graphical interface for processing *trace-cmd* trace.dat reports.
- ▶ Can also setup and acquire data using *trace-cmd*.
- ▶ Displays CPU and tasks as different colors along with the recorded events.
- ▶ Useful when a deep analysis is required for a specific bug.







Profiling a system from userspace to kernel space

- ▶ Profiling with ftrace, uprobes and kernelshark
- ▶ Profiling with perf



LTTng



- ▶ LTTng is an open source tracing framework for Linux maintained by the [EfficiOS](#) company.
- ▶ LTTng allows understanding the interactions between the kernel and applications (C, C++, Java, Python).
 - Also expose a `/dev/lttng-logger` that can be used from any application.
- ▶ Tracepoints are associated with a payload (data).
- ▶ LTTng is focused on low-overhead tracing.
- ▶ Uses the Common Trace Format (so traces are readable with other software like babeltrace or trace-compass)





Tracepoints with LTTng

- ▶ LTTng works with a session daemon that receive all events from kernel and userspace LTTng tracing components.
- ▶ LTTng can use and trace the following instrumentation points:
 - LTTng kernel tracepoints
 - kprobes and kretprobes
 - Linux kernel system calls
 - Linux user space probe
 - User space LTTng tracepoints



Creating userspace tracepoints with LTTng

- ▶ New userspace tracepoints can be defined using LTTng.
- ▶ Tracepoints have multiple characteristics:
 - A provider namespace
 - A name identifying the tracepoint
 - Parameters of various types (int, char *, etc)
 - Fields describing how to display the tracepoint parameters (decimal, hexadecimal, etc) (see [LTTng-ust](#) manpage for types)
- ▶ Developers must perform multiple operations to use UST tracepoint: write a tracepoint provider (.h), write a tracepoint package (.c), build the package, call the tracepoint in the traced application, and finally build the application, linked with lttng-ust library and the package provider.
- ▶ LTTng provides the `lttng-gen-tp` to ease all those steps, allowing to only write a template (.tp) file.



Defining a LTTng tracepoint (1/2)

► Tracepoint template (hello_world-tp.tp):

```
LTTNG_UST_TRACEPOINT_EVENT(  
    // Tracepoint provider name  
    hello_world,  
  
    // Tracepoint/event name  
    my_first_tracepoint,  
  
    // Tracepoint arguments (input)  
    LTTNG_UST_TP_ARGS(  
        char *, text  
    ),  
  
    // Tracepoint/event fields (output)  
    LTTNG_UST_TP_FIELDS(  
        lttng_ust_field_string(message, text)  
    )  
)
```

► lttng-gen-tp will take this template file and generate/build all needed files (.h, .c and .o files)



Defining a LTTng tracepoint (2/2)

▶ Build tracepoint provider:

```
$ lttng-gen-tp hello_world-tp.tp
```

▶ Tracepoint usage (hello_world.c):

```
#include <stdio.h>
#include "hello-tp.h"

int main(int argc, char *argv[])
{
    lttng_ust_tracepoint(hello_world, my_first_tracepoint, "hi there!");
    return 0;
}
```

▶ Compilation:

```
$ gcc hello_world.c hello_world-tp.o -llttng-ust -o hello_world
```



Using LTTng

```
$ lttng create my-tracing-session --output=./my_traces
$ lttng list --kernel
$ lttng list --userspace
$ lttng enable-event --userspace hello_world:my_first_tracepoint
$ lttng enable-event --kernel --syscall open,close,write
$ lttng start
$ /* Run your application or do something */
$ lttng destroy
$ babeltrace2 ./my_traces
```

- ▶ You can also use [trace-compass](#) to display the traces in a GUI



Remote tracing with LTTng

- ▶ LTTng allows to record traces over the network.
- ▶ Useful for embedded systems with limited storage capabilities.
- ▶ On the remote computer, run `lttng-relayd` command

```
$ lttng-relayd --output=${PWD}/traces
```

- ▶ Then on the target, at session creation, use the `--set-url`

```
$ lttng create my-session --set-url=net://remote-system
```

- ▶ Traces will then be recorded directly on the remote computer.



eBPF




The ancestor: Berkeley Packet filter

- ▶ BPF stands for Berkeley Packet Filter and was initially used for network packet filtering
- ▶ BPF is implemented and used in Linux to perform Linux Socket Filtering (see [networking/filter](#))
- ▶ tcpdump and Wireshark heavily rely on BPF (through libpcap) for packet capture



BPF in libpcap: setup


- ▶ tcpdump passes the capture filter string from the user to libpcap
- ▶ libpcap translates the capture filter into a binary program
 - This program uses the instruction set of an abstract machine (the “BPF instruction set”)
- ▶ libpcap sends the binary program to the kernel via the `setsockopt()` syscall





BPF in libpcap: capture

- ▶ The kernel implements the BPF “virtual machine”
- ▶ The BPF virtual machine executes the program for every packet
- ▶ The program inspects the packet data and returns a non-zero value if the packet must be captured
- ▶ If the return value is non-zero, the packet is captured in addition to regular packet processing





eBPF (1/2)

- ▶ eBPF is a new framework allowing to run small user programs directly in the kernel, in a safe and efficient way. It has been added in kernel 3.18 but it is still evolving and receiving updates frequently.
- ▶ eBPF programs can capture and expose kernel data to userspace, and also alter kernel behavior based on some user-defined rules.
- ▶ eBPF is event-driven: an eBPF program is triggered and executed on a specific kernel event
- ▶ A major benefit from eBPF is the possibility to reprogram the kernel behavior, without performing kernel development:
 - no risk of crashing the kernel because of bugs
 - faster development cycles to get a new feature ready




Image credits: <https://ebpf.io/>



- ▶ The most notable eBPF features are:
 - A new instruction set, interpreter and verifier
 - A wide variety of "attach" locations, allowing to hook programs almost anywhere in the kernel
 - dedicated data structures called "maps", to exchange data between multiple eBPF programs or between programs and userspace
 - A dedicated `bpf()` syscall to manipulate eBPF programs and data
 - plenty of (kernel) helper functions accessible from eBPF programs.



eBPF program lifecycle





Kernel configuration for eBPF

- ▶ `CONFIG_NET` to enable eBPF subsystem
- ▶ `CONFIG_BPF_SYSCALL` to enable the `bpf()` syscall
- ▶ `CONFIG_BPF_JIT` to enable JIT on programs and so increase performance
- ▶ `CONFIG_BPF_JIT_ALWAYS_ON` to force JIT
- ▶ `CONFIG_BPF_UNPRIV_DEFAULT_OFF=n` in **development** to allow eBPF usage without root
- ▶ You may then want to enable more general features to "unlock" specific hooking locations:
 - `CONFIG_KPROBES` to allow hooking programs on kprobes
 - `CONFIG_TRACING` to allow hooking programs on kernel tracepoints
 - `CONFIG_NET_CLS_BPF` to write packets classifiers
 - `CONFIG_CGROUP_BPF` to attach programs on cgroups hooks



- ▶ eBPF is a "virtual" ISA, defining its own set of instructions: load and store instructions, arithmetic instructions, jump instructions, etc
- ▶ It also defines a set of 10 64-bits wide registers as well as a calling convention:
 - R0: return value from functions and BPF program
 - R1, R2, R3, R4, R5: function arguments
 - R6, R7, R8, R9: callee-saved registers
 - R10: stack pointer

```
; bpf_printk("Hello %s\n", "World");
0:  r1 = 0x0 11
2:  r2 = 0xa
3:  r3 = 0x0 11
5:  call 0x6
; return 0;
6:  r0 = 0x0
7:  exit
```




The eBPF verifier

- ▶ When loaded into the kernel, a program must first be validated by the eBPF verifier.
- ▶ The verifier is a complex piece of software which checks eBPF programs against a set of rules to ensure that running those may not compromise the whole kernel. For example:
 - a program must always return and so not contain paths which could make them "infinite" (e.g: no infinite loop)
 - a program must make sure that a pointer is valid before dereferencing it
 - a program cannot access arbitrary memory addresses, it must use passed context and available helpers
- ▶ If a program violates one of the verifier rules, it will be rejected.
- ▶ Despite the presence of the verifier, you still need to be careful when writing programs! eBPF programs run with preemption enabled (but CPU migration disabled), so they can still suffer from concurrency issues
 - There are mechanisms and helpers to avoid those issues, like per-CPU maps types.



Program types and attach points

- ▶ There are different categories of hooks to which a program can be attached:
 - an arbitrary kprobe
 - a kernel-defined static tracepoint
 - a specific perf event
 - throughout the network stack
 - an arbitrary uprobe
 - and a lot more, see [bpf_attach_type](#)
- ▶ A specific attach-point type can only be hooked with a set of specific program types, see [bpf_prog_type](#) and [bpf/libbpf/program_types](#).
- ▶ The program type then defines the data passed to an eBPF program as input when it is invoked. For example:
 - A `BPF_PROG_TYPE_TRACEPOINT` program will receive a structure containing all data returned to userspace by the targeted tracepoint.
 - A `BPF_PROG_TYPE_SCHED_CLS` program (used to implement packet classifiers) will receive a [struct __sk_buff](#), the kernel representation of a socket buffer.
 - You can learn about the context passed to any program type by checking [include/linux/bpf_types.h](#)



- ▶ eBPF programs exchange data with userspace or other programs through maps of different natures:
 - `BPF_MAP_TYPE_ARRAY`: generic array storage. Can be differentiated per CPU
 - `BPF_MAP_TYPE_HASH`: a storage composed of key-value pairs. Keys can be of different types: `__u32`, a device type, an IP address...
 - `BPF_MAP_TYPE_QUEUE`: a FIFO-type queue
 - `BPF_MAP_TYPE_CGROUP_STORAGE`: a specific hash map keyed by a cgroup id. There are other types of maps specific to other object types (inodes, tasks, sockets, etc)
 - etc...
- ▶ For basic data, it is easier and more efficient to directly use eBPF global variables (no syscalls involved, contrary to maps)



The `bpf()` syscall

- ▶ The kernel exposes a `bpf()` syscall to allow interacting with the eBPF subsystem
- ▶ The syscall takes a set of subcommands, and depending on the subcommand, some specific data:
 - `BPF_PROG_LOAD` to load a bpf program
 - `BPF_MAP_CREATE` to allocate maps to be used by a program
 - `BPF_MAP_LOOKUP_ELEM` to search for an entry in a map
 - `BPF_MAP_UPDATE_ELEM` to update an entry in a map
 - etc
- ▶ The syscall works with file descriptors pointing to eBPF resources. Those resources (program, maps, links, etc) remain valid while there is at least one program holding a valid file descriptor to it. Those are automatically cleaned once there are no user left.
- ▶ For more details, see `man 2 bpf`



Writing eBPF programs

- ▶ eBPF programs can either be written directly in raw eBPF assembly or in higher level languages (e.g: C or rust), and are compiled using the clang compiler.
- ▶ The kernel provides some helpers that can be called from an eBPF program:
 - `bpf_trace_printk` Emits a log to the trace buffer
 - `bpf_map_{lookup,update,delete}_elem` Manipulates maps
 - `bpf_probe_{read,write}[_user]` Safely read/write data from/to kernel or userspace
 - `bpf_get_current_pid_tgid` Returns current Process ID and Thread group ID
 - `bpf_get_current_uid_gid` Returns current User ID and Group ID
 - `bpf_get_current_comm` Returns the name of the executable running in the current task
 - `bpf_get_current_task` Returns the current `struct task_struct`
 - Many other helpers are available, see [man 7 bpf-helpers](#)
- ▶ Kernel also exposes kfuncs (see [bpf/kfuncs](#)), but contrary to bpf-helpers, those do not belong to the kernel stable interface.



Manipulating eBPF program

- ▶ There are different ways to build, load and manipulate eBPF programs:
 - One way is to write an eBPF program, build it with clang, and then load it, attach it and read data from it with bare `bpf()` calls in a custom userspace program
 - One can also use `bpftool` on the built ebpf program to manipulate it (load, attach, read maps, etc), without writing any userspace tool
 - Or we can write our own eBPF tool thanks to some intermediate libraries which handle most of the hard work, like `libbpf`
 - We can also use specialized frameworks like BCC or `bpftool` to really get all operations (bpf program build included) handled



- ▶ BPF Compiler Collection (BCC) is (as its name suggests) a collection of BPF based tools.
- ▶ BCC provides a large number of ready-to-use tools written in BPF.
- ▶ Also provides an interface to write, load and hook BPF programs more easily than using "raw" BPF language.
- ▶ Available on a large number of architectures and distributions (but not packaged in Buildroot)
 - On debian, when installed, all tools are named `<tool>-bpfcc`.
- ▶ BCC requires a kernel version ≥ 4.1 .
- ▶ BCC evolves quickly, many distributions have old versions: you may need to compile from the latest sources




Image credits:

<https://github.com/iovisor/bcc>



BCC tools

Linux bcc/BPF Tracing Tools



<https://github.com/iovisor/bcc#tools> 2019

Image credits: <https://www.brendangregg.com/ebpf.html>



BCC Tools example

- ▶ `profile.py` is a CPU profiler allowing to capture stack traces of current execution. Its output can be used for flamegraph generation:

```
$ git clone https://github.com/brendangregg/FlameGraph.git
$ profile.py -df -F 99 10 | ./FlameGraph/flamegraph.pl > flamegraph.svg
```

- ▶ `tcpconnect.py` script displays all new TCP connections live

```
$ tcpconnect
PID      COMM      IP SADDR      DADDR      DPORT
220321  ssh        6  ::1         ::1         22
220321  ssh        4  127.0.0.1   127.0.0.1   22
17676   Chrome_Child 6  2a01:cb15:81e4:8100:37cf:d45b:d87d:d97d 2606:50c0:8003::154 443
[...]
```

- ▶ And much more to discover at <https://github.com/iovisor/bcc>



Using BCC with python

- ▶ BCC exposes a `bcc` module, and especially a `BPF` class
- ▶ eBPF programs are written in C and stored either in external files or directly in a python string.
- ▶ When an instance of the `BPF` class is created and fed with the program (either as string or file), it automatically builds, loads, and possibly attaches the program
- ▶ There are multiple ways to attach a program:
 - By using a proper program name prefix, depending on the targeted attach point (and so the attach step is performed automatically)
 - By explicitly calling the relevant attach method on the `BPF` instance created earlier



Using BCC with python

- ▶ Hook with a *kprobe* on the `clone()` system call and display "Hello, World!" each time it is called

```
#!/usr/bin/env python3

from bcc import BPF

# define BPF program
prog = '''
int hello(void *ctx) {
    bpf_trace_printk("Hello, World!\\n");
    return 0;
}
'''

# load BPF program
b = BPF(text=prog)
b.attach_kprobe(event=b.get_syscall_fnname("clone"), fn_name="hello")
```



- ▶ Creating custom tracing tools with BCC framework



- ▶ Instead of using a high level framework like BCC, one can use libbpf to build custom tools with finer control over every aspect of the program.
- ▶ libbpf is a C-based library that aims to ease eBPF programming thanks to the following features:
 - userspace APIs to handle open/load/attach/teardown of bpf programs
 - userspace APIs to interact with attached programs
 - eBPF APIs to ease eBPF program writing
- ▶ Packaged in many distributions and build systems (e.g.: Buildroot)
- ▶ Learn more at <https://libbpf.readthedocs.io/en/latest/>



eBPF programming with libbpf (1/2)

my_prog.bpf.c

```
#include <linux/bpf.h>
#include <bpf/bpf_helpers.h>
#include <bpf/bpf_tracing.h>

#define TASK_COMM_LEN 16
struct {
    __uint(type, BPF_MAP_TYPE_ARRAY);
    __type(key, __u32);
    __type(value, __u64);
    __uint(max_entries, 1);
} counter_map SEC(".maps");

struct sched_switch_args {
    unsigned long long pad;
    char prev_comm[TASK_COMM_LEN];
    int prev_pid;
    int prev_prio;
    long long prev_state;
    char next_comm[TASK_COMM_LEN];
    int next_pid;
    int next_prio;
};
```

- ▶ The fields to define in the *_args structure are obtained from the event description in /sys/kernel/tracing/events (see [this example](#))



eBPF programming with libbpf (2/2)

my_prog.bpf.c

```
SEC("tracepoint/sched/sched_switch")
int sched_tracer(struct sched_switch_args *ctx)
{
    __u32 key = 0;
    __u64 *counter;
    char *file;

    char fmt[] = "Old task was %s, new task is %s\n";
    bpf_trace_printk(fmt, sizeof(fmt), ctx->prev_comm, ctx->next_comm);

    counter = bpf_map_lookup_elem(&counter_map, &key);
    if(counter) {
        *counter += 1;
        bpf_map_update_elem(&counter_map, &key, counter, 0);
    }

    return 0;
}

char LICENSE[] SEC("license") = "Dual BSD/GPL";
```



Building eBPF programs

- ▶ An eBPF program written in C can be built into a loadable object thanks to clang:

```
$ clang -target bpf -O2 -g -c my_prog.bpf.c -o my_prog.bpf.o
```

- The `-g` option allows to add debug information as well as BTF information
- ▶ GCC can be used too with recent versions
 - the toolchain can be installed with the `gcc-bpf` package in Debian/Ubuntu
 - it exposes the `bpf-unknown-none` target
- ▶ To easily manipulate this program with a userspace program based on libbpf, we need "skeleton" APIs, which can be generated with `bpftool`



- ▶ `bpftool` is a command line tool allowing to interact with bpf object files and the kernel to manipulate bpf programs:
 - Load programs into the kernel
 - List loaded programs
 - Dump program instructions, either as BPF code or JIT code
 - List loaded maps
 - Dump map content
 - Attach programs to hooks (so they can run)
 - etc
- ▶ You may need to mount the bpf filesystem to be able to pin a program (needed to keep a program loaded after bpftool has finished running):

```
$ mount -t bpf none /sys/fs/bpf
```



▶ List loaded programs

```
$ bpftool prog
348: tracepoint name sched_tracer tag 3051de4551f07909 gpl
loaded_at 2024-08-06T15:43:11+0200 uid 0
xlated 376B jited 215B memlock 4096B map_ids 146,148
btf_id 545
```

▶ Load and attach a program

```
$ mkdir /sys/fs/bpf/myprog
$ bpftool prog loadall trace_execve.bpf.o /sys/fs/bpf/myprog autoattach
```

▶ Unload a program

```
$ rm -rf /sys/fs/bpf/myprog
```



► Dump a loaded program

```
$ bpftool prog dump xlated id 348
int sched_tracer(struct sched_switch_args * ctx):
; int sched_tracer(struct sched_switch_args *ctx)
  0: (bf) r4 = r1
  1: (b7) r1 = 0
; __u32 key = 0;
  2: (63) *(u32 *)(r10 -4) = r1
; char fmt[] = "Old task was %s, new task is %s\n";
  3: (73) *(u8 *)(r10 -8) = r1
  4: (18) r1 = 0xa7325207369206b
  6: (7b) *(u64 *)(r10 -16) = r1
  7: (18) r1 = 0x7361742077656e20
[...]
```

► Dump eBPF program logs

```
$ bpftool prog tracelog
kworker/u80:0-11 [013] d..41 1796.003605: bpf_trace_printk: Old task was kworker/u80:0, new task is swapper/13
<idle>-0 [013] d..41 1796.003609: bpf_trace_printk: Old task was swapper/13, new task is kworker/u80:0
sudo-18640 [010] d..41 1796.003613: bpf_trace_printk: Old task was sudo, new task is swapper/10
<idle>-0 [010] d..41 1796.003617: bpf_trace_printk: Old task was swapper/10, new task is sudo
[...]
```



► List created maps

```
$ bpftool map
80: array name counter_map flags 0x0
    key 4B value 8B max_entries 1 memlock 256B
    btf_id 421
82: array name .rodata.str1.1 flags 0x80
    key 4B value 33B max_entries 1 memlock 288B
    frozen
96: array name libbpf_global flags 0x0
    key 4B value 32B max_entries 1 memlock 280B
[...]
```

► Show a map content

```
$ sudo bpftool map dump id 80
[{"key": 0,
  "value": 4877514
}]
```



- ▶ Generate libbpf APIs to manipulate a program

```
$ bpftool gen skeleton trace_sched_switch.bpf.o name trace_sched_switch \  
> trace_sched_switch.skel.h
```

- ▶ We can then write our userspace program and benefit from high level APIs to manipulate our eBPF program:
 - instantiation of a global context object which will have references to all of our programs, maps, links, etc
 - loading/attaching/unloading of our programs
 - eBPF program directly embedded in the generated header as a byte array



Userspace code with libbpf

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include "trace_sched_switch.skel.h"

int main(int argc, char *argv[])
{
    struct trace_sched_switch *skel;
    int key = 0;
    long counter = 0;

    skel = trace_sched_switch__open_and_load();
    if(!skel)
        exit(EXIT_FAILURE);
    if (trace_sched_switch__attach(skel)) {
        trace_sched_switch__destroy(skel);
        exit(EXIT_FAILURE);
    }

    while(true) {
        bpf_map__lookup_elem(skel->maps.counter_map, &key, sizeof(key), &counter, sizeof(counter), 0);
        fprintf(stderr, "Scheduling switch count: %d\n", counter);
        sleep(1);
    }

    return 0;
}
```



eBPF programs portability (1/2)

- ▶ Kernel internals, contrary to userspace APIs, do not expose stable APIs. This means that an eBPF program manipulating some kernel data may not work with another kernel version
- ▶ The CO-RE (Compile Once - Run Everywhere) approach aims to solve this issue and make programs portable between **kernel versions**. It relies on the following features:
 - your kernel must be built with `CONFIG_DEBUG_INFO_BTF=y` to have BTF data embedded. BTF is a format similar to dwarf which encodes data layout and function signatures in an efficient way.
 - your eBPF compiler must be able to emit BTF relocations (both clang and GCC are capable of this on recent versions, with the `-g` argument)
 - you need a BPF loader capable of processing BPF programs based on BTF data and adjust accordingly data access: `libbpf` is the de-facto standard bpf loader
 - you then need eBPF APIs to read/write to CO-RE relocatable variables. `libbpf` provides such helpers, like `bpf_core_read`
- ▶ To learn more, take a look at [Andrii Nakryiko's CO-RE guide](#)



eBPF programs portability (2/2)

- ▶ Despite CO-RE, you may still face different constraints on different kernel versions, because of major features introduction or change, since the eBPF subsystem keeps receiving frequent updates:
 - eBPF tail calls (which allow a program to call a function) have been added in version 4.2, and allow to call another program only since version 5.10
 - eBPF spin locks have been added in version 5.1 to prevent concurrent access to maps shared between CPUs.
 - Different attach types keep being added, but possibly on different kernel versions when it depends on the architecture: fentry/fexit attach points have been added in kernel 5.5 for x86 but in 6.0 for arm32.
 - Any kind of loop (even bounded) was forbidden until version 5.3
 - CAP_BPF capability, allowing a process to perform eBPF tasks, has been added in version 5.8




eBPF for tracing/profiling

- ▶ eBPF is a very powerful framework to spy on kernel internals: thanks to the wide variety of attach point, you can expose almost any kernel code path and data.
- ▶ In the meantime, eBPF programs remain isolated from kernel code, which makes it safe (compared to kernel development) and easy to use.
- ▶ Thanks to the in-kernel interpreter and optimizations like JIT compilation, eBPF is very well suited for tracing or profiling with low overhead, even in production environments, while being very flexible.
- ▶ This is why eBPF adoption level keeps growing for debugging, tracing and profiling in the Linux ecosystem. As a few examples, we find eBPF usage in:
 - tracing frameworks like [BCC](#) and [bpftrace](#)
 - network infrastructure components, like [Cilium](#) or [Calico](#)
 - network packet tracers, like [pwru](#) or [dropwatch](#)
 - And many more, check [ebpf.io](#) for more examples



- ▶ libbpf-bootstrap: <https://github.com/libbpf/libbpf-bootstrap>
- ▶ A Beginner's Guide to eBPF Programming - Liz Rice, 2020
 - Video: <https://www.youtube.com/watch?v=lrSEXTfS-iQ>
 - Resources: <https://github.com/lizrice/ebpf-beginners>





Porting our custom tracing tool for embedded use case

- ▶ Converting a BCC script to libbpf
- ▶ Bringing advanced features to the tool



Choosing the right tool




Choosing the right tool

- ▶ Before starting to profile or trace, one should know which type of tool to use.
- ▶ This choice is guided by the level of profiling
- ▶ Often start by analyzing/optimizing the application level using application tracing/profiling tools (valgrind, perf, etc).
- ▶ Then analyze user space + kernel performance
- ▶ Finally, trace or profile the whole system if the performance problems happens only when running under a loaded system.
 - For "constant" load problems, snapshot tools works fine.
 - For sporadic problems, record traces and analyze them.
- ▶ If you happen to have a complex setup that you often have to bring up, it is likely a sign that you want to ease this setup with some custom tooling: scripting, custom traces, eBPF, etc



Kernel Debugging

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Preventing bugs



Static code analysis

- ▶ Static analysis can be run with the *sparse* tool
- ▶ *sparse* works with annotation and can detect various errors at compile time
 - Locking issues (unbalanced locking)
 - Address space issues, such as accessing user space pointer directly
- ▶ Analysis can be run using `make C=2` to run only on files that are recompiled
- ▶ Or with `make C=1` to run on all files
- ▶ Example of an unbalanced locking scheme:

```
rzn1_a5psw.c:81:13: warning: context imbalance in 'a5psw_reg_rmw' - wrong count  
at exit
```

SP[^]ARSE



Good practices in kernel development (1/2)

- ▶ When writing driver code, never expect the user to provide correct values. Always check these values.
- ▶ Use the `WARN_ON()` macro if you want to display a stacktrace when a specific condition did happen.
 - `dump_stack()` can also be used during debugging to show the current call stack.

```
static bool check_flags(u32 flags)
{
    if (WARN_ON(flags & STATE_INVALID))
        return -EINVAL;
    return 0;
}
```



Good practices in kernel development (2/2)

- ▶ If the values can be checked at compile time (configuration input, `sizeof`, structure fields), use the `BUILD_BUG_ON()` macro to ensure the condition is true.

```
BUILD_BUG_ON(sizeof(ctx->__reserved) != sizeof(reserved));
```

- ▶ If during compilation you have some warnings about unused variables/parameters, they must be fixed.
- ▶ Apply `checkpatch.pl --strict` when possible which might find some potential problems in your code.



Linux Kernel Debugging



Linux Kernel Debugging

- ▶ The Linux Kernel features multiple tools to ease kernel debugging:
 - A dedicated logging framework
 - A standard way to dump low level crash messages
 - Multiple runtime checkers to check for different kind of issues: memory issues, locking mistakes, undefined behaviors, etc.
 - Interactive or post-mortem debugging
- ▶ Many of those features need to be explicitly enabled in the kernel menuconfig, those are grouped in the `Kernel hacking -> Kernel debugging` menuconfig entry.
 - `CONFIG_DEBUG_KERNEL` should be set to "y" to enable other debug options.



Debugging using messages



Debugging/tracing using logs 1/4

▶ Good old `printk()`!

- Works in all contexts
- Can specify a log level ranging from 0 (emergency) to 7 (debug)
- Be careful of the delays introduced when logs are spitted out on a serial console at 115200 bauds
 - A `*_ratelimited()` version exists to limit the amount of print if called too often
- Not recommended for upstream contributions

Example:

```
printk("in probe\n");
```

Here's what you get in the kernel log:

```
[ 1.878382] in probe
```

All other logging facilities are based on it.



Debugging/tracing using logs 2/4

- ▶ The `pr_*`() family of functions
 - They include the log level in the name:
`pr_emerg()`, `pr_alert()`, `pr_crit()`, `pr_err()`, `pr_warn()`, `pr_notice()`,
`pr_info()`, `pr_cont()` and the special `pr_debug()` (see next pages)
 - They allow setting a manual prefix (eg. eases grepping):
`#define pr_fmt(fmt) "foo: " fmt`
- ▶ Also defined in `include/linux/printk.h`

Example:

```
pr_info("in probe\n");
```

Here's what you get in the kernel log:

```
[    1.878382] in probe
```

or similarly with a manual format:

```
[    1.878382] foo: in probe
```



Debugging/tracing using logs 3/4

► The `dev_*()` family of functions

- They include a formatted standard prefix with the device name:
`dev_emerg()`, `dev_alert()`, `dev_crit()`, `dev_err()`, `dev_warn()`, `dev_notice()`,
`dev_info()` and the special `dev_dbg()` (see next pages)
- They additionally take a pointer to `struct device` as first argument
- Defined in `include/linux/dev_printk.h`
- To be used in device drivers

Example:

```
dev_info(&pdev->dev, "in probe\n");
```

Here's what you get in the kernel log:

```
[    1.878382] serial 48024000.serial: in probe
[    1.884873] serial 481a8000.serial: in probe
```




Debugging/tracing using logs 4/4

- ▶ The kernel defines many more format specifiers than the standard `printf()` existing ones.
 - `%p`: Display the hashed value of pointer by default.
 - `%px`: Always display the address of a pointer (use carefully on non-sensitive addresses).
 - `%pK`: Display hashed pointer value, zeros or the pointer address depending on `kptr_restrict` sysctl value.
 - `%pOF`: Device-tree node format specifier.
 - `%pr`: Resource structure format specifier.
 - `%pa`: Physical address display (work on all architectures 32/64 bits)
 - `%pe`: Error pointer (displays the string corresponding to the error number)
- ▶ See [core-api/printk-formats](#) for an exhaustive list of format specifiers
- ▶ Also features a helper to dump entire buffers with a `hexdump` like display: [print_hex_dump\(\)](#)



pr_debug() and dev_dbg()

- ▶ When the driver is compiled with `DEBUG` defined, all these messages are compiled and printed at the debug level. `DEBUG` can be defined by `#define DEBUG` at the beginning of the driver, or using `ccflags-$(CONFIG_DRIVER) += -DDEBUG` in the `Makefile`
- ▶ When the kernel is compiled with `CONFIG_DYNAMIC_DEBUG`, then these messages can dynamically be enabled on a per-file, per-module or per-message basis, by writing commands to `/proc/dynamic_debug/control`. Note that messages are not enabled by default.
 - Details in [admin-guide/dynamic-debug-howto](#)
 - Very powerful feature to only get the debug messages you're interested in.
- ▶ When neither `DEBUG` nor `CONFIG_DYNAMIC_DEBUG` are used, these messages are not compiled in.



pr_debug() and dev_dbg() usage

- ▶ Debug prints can be enabled using the `/proc/dynamic_debug/control` file.
 - `cat /proc/dynamic_debug/control` will display all lines that can be enabled in the kernel
 - Example: `init/main.c:1427 [main]run_init_process =p " \%s\012"`
- ▶ A syntax allows to enable individual print using lines, files or modules
 - `echo "file drivers/pinctrl/core.c +p" > /proc/dynamic_debug/control` will enable all debug prints in `drivers/pinctrl/core.c`
 - `echo "module pciehp +p" > /proc/dynamic_debug/control` will enable the debug print located in the `pciehp` module
 - `echo "file init/main.c line 1427 +p" > /proc/dynamic_debug/control` will enable the debug print located at line 1247 of file `init/main.c`
 - Replace `+p` with `-p` to disable the debug print



Debug logs troubleshooting

- ▶ When using dynamic debug, make sure that your debug call is enabled: it must be visible in `control` file in debugfs **and** be activated (`=p`)
- ▶ Is your log output only in the kernel log buffer?
 - You can see it thanks to `dmesg`
 - You can lower the `loglevel` to output it to the console directly
 - You can also set `ignore_loglevel` in the kernel command line to force all kernel logs to console
- ▶ If you are working on an out-of-tree module, you may prefer to define `DEBUG` in your module source or Makefile instead of using dynamic debug
- ▶ If configuration is done through the kernel command line, is it properly interpreted?
 - Starting from 5.14, kernel will let you know about faulty command line:
Unknown kernel command line parameters `loglevel`, will be passed to user space.
 - You may need to take care of special characters escaping (e.g: quotes)
- ▶ Be aware that a few subsystems bring their own logging infrastructure, with specific configuration/controls, eg: `drm.debug=0x1ff`



Kernel early debug

- ▶ When booting, the kernel sometimes crashes even before displaying the system messages
- ▶ On ARM, if your kernel doesn't boot or hangs without any message, you can activate early debugging options
 - `CONFIG_DEBUG_LL=y` to enable ARM early serial output capabilities
 - `CONFIG_EARLY_PRINTK=y` will allow printk to output the prints earlier
- ▶ `earlyprintk` command line parameter should be given to enable early printk output




Kernel crashes and oops



Kernel crashes

- ▶ The kernel is not immune to crash, many errors can be done and lead to crashes
 - Memory access error (NULL pointer, out of bounds access, etc)
 - Voluntarily panicking on error detection (using `panic()`)
 - Kernel incorrect execution mode (sleeping in atomic context)
 - Deadlocks detected by the kernel (Soft lockup/locking problem)
- ▶ On error, the kernel will display a message on the console that is called a "Kernel oops"



Icon by Peter van Driel, TheNounProject.com



Kernel oops (1/2)

- ▶ The content of this message depends on the architecture that is used.
- ▶ Almost all architectures display at least the following information:
 - CPU state when the oops happened
 - Registers content with potential interpretation
 - Backtrace of function calls that led to the crash
 - Stack content (last X bytes)
- ▶ Depending on the architecture, the crash location can be identified using the content of the PC registers (sometimes named IP, EIP, etc).
- ▶ To have a meaningful backtrace with symbol names use `CONFIG_KALLSYMS=y` which will embed the symbol names in the kernel image.



Kernel oops (2/2)

- ▶ Symbols are displayed in the backtrace using the following format:
 - `<symbol_name>+<hex_offset>/<symbol_size>`
- ▶ If the oops is not critical (taken in process context), then the kernel will kill process and continue its execution
 - The kernel stability might be compromised!
- ▶ Tasks that are taking too much time to execute and that are hung can also generate an oops (`CONFIG_DETECT_HUNG_TASK`)
- ▶ If KGDB support is present and configured, on oops, the kernel will switch to KGDB mode.



Oops example (1/2)

```
1.635900] 8<--- cut here ---
1.639034] Unable to handle kernel NULL pointer dereference at virtual address 00000050
1.647304] [00000050] *pgd=00000000
1.650959] Internal error: Oops: 5 [#1] SMP ARM
1.655635] Modules linked in:
1.658736] CPU: 0 PID: 28 Comm: kworker/u4:1 Not tainted 5.10.0-rc6-01642-g59dcb02d20b #40
1.667261] Hardware name: Generic DT based system
1.672164] Workqueue: events unbound_deferred_probe_work_func
1.678029] PC is at mlic_create+0xac/0x134
1.682280] LR is at __raw_spin_unlock_irqrestore+0x2c/0x34
1.687841] pc : [<804ba6e4>] lr : [<806811f0>] psr: 20000053
1.694172] sp : 90a99cf8 ip : 90a99cc8 fp : 90a99d1c
1.699453] r10: 8089339e r9 : 808a5a18 r8 : 8fdb4284
1.704733] r7 : 82173000 r6 : 8fdb47d8 r5 : 82173010 r4 : 8fdb3ad8
1.711327] r3 : 00000000 r2 : 00000000 r1 : a0000053 r0 : 8fdb3ad8
1.717921] Flags: nZCv IRQs on FIQs off Mode SVC_32 ISA ARM Segment none
1.725220] Control: 10c53b7d Table: 8000406a DAC: 00000051
1.731022] Register r0 information: non-slab/vmalloc memory
1.736754] Register r1 information: non-paged memory
1.741867] Register r2 information: NULL pointer
1.746629] Register r3 information: NULL pointer
1.751389] Register r4 information: non-slab/vmalloc memory
1.757113] Register r5 information: slab kmallo-512 start 82173000 pointer offset 16 size 512
1.765943] Register r6 information: non-slab/vmalloc memory
1.771668] Register r7 information: slab kmallo-512 start 82173000 pointer offset 0 size 512
1.780401] Register r8 information: non-slab/vmalloc memory
1.786125] Register r9 information: non-slab/vmalloc memory
1.791848] Register r10 information: non-slab/vmalloc memory
1.797661] Register r11 information: 2-page vmalloc region starting at 0x90a98000 allocated at kernel_clone+0xb4/0x2a4
1.808567] Register r12 information: 2-page vmalloc region starting at 0x90a98000 allocated at kernel_clone+0xb4/0x2a4
1.819468] Process kworker/u4:1 (pid: 28, stack limit = 0x(ptrval))
1.825889] Stack: (0x90a99cf8 to 0x90a9a000)
1.830304] 9ce0: 00000005 ddf3fe37
1.838570] 9d00: 8259c840 8259c840 82173010 8fdb47d8 90a99dac 90a99d20 804bb7b8 804ba644
1.846834] 9d20: 00000000 90a99d34 00000002 8fdb46bc 00000000 8fdb4284 00000000 804f1c80
1.855099] 9d40: ffffffff 00000001 8089ad70 ffffffff 8fdb43f8 ff8fb728 ff8fb728 ddf3fe37
1.863364] 9d60: 82173010 82173010 81383ea0 82173010 00000002 8219cac0 90a99d9c ddf3fe37
1.871627] 9d80: 80458940 00000000 82173010 81383ea0 82173010 00000002 8219cac0 8200f00d
1.879892] 9da0: 90a99dcc 90a99db0 80449f44 804bb664 00000000 82173010 81383ea0 82173010
1.888155] 9dc0: 90a99dec 90a99dd0 80447cb8 80449ee8 82173010 81383ea0 81383ea0 82173010
1.896420] 9de0: 90a99e04 90a99df0 80447ef0 80447b9c 815b034c 815b0350 90a99e2c 90a99e08
1.904684] 9e00: 80447f54 80447e28 81383ea0 90a99e84 82173010 00000001 00000000 8219cac0
1.912948] 9e20: 90a99e4c 90a99e30 80448320 80447f10 00000000 90a99e84 8044827c 00000001
1.921213] 9e40: 90a99e7c 90a99e50 80445dac 80448288 82173010 82098e6c 82575ab8 ddf3fe37
1.929478] 9e60: 90a99e7c 82173010 8137c268 82173054 90a99eac 90a99e80 80448158 80445d00
1.937742] 9e80: 90a99eac 82173010 00000001 ddf3fe37 82173010 8137c268 82173010 8200f000
1.946006] 9ea0: 80447f54 80447e28 81383ea0 90a99e84 82173010 82098e6c 82575ab8 ddf3fe37
```

Cause and
generic information

Register content
and CPU state

Stack content



Oops example (2/2)

```
2.036894] Backtrace:
2.039380] miic create from a5psw_probe+0x160/0x3e0
2.044530] r6:8fdb47d8 r5:82173010 r4:8259c840
2.049194] a5psw_probe from platform_probe+0x68/0xb8
2.054432] r10:8200f00d r9:8219cac0 r8:00000002 r7:82173010 r6:81383ea0 r5:82173010
2.062332] r4:00000000
2.064896] platform_probe from really_probe+0x128/0x28c
2.070383] r7:82173010 r6:81383ea0 r5:82173010 r4:00000000
2.076095] really_probe from __driver_probe_device+0xd4/0xe8
2.082018] r7:82173010 r6:81383ea0 r5:81383ea0 r4:82173010
2.087730] __driver_probe_device from driver_probe_device+0x50/0xd0
2.094260] r5:815b0350 r4:815b034c
2.097874] driver_probe_device from __device_attach_driver+0xa4/0xc4
2.104499] r9:8219cac0 r8:00000000 r7:00000001 r6:82173010 r5:90a99e84 r4:81383ea0
2.112309] __device_attach_driver from bus_for_each_drv+0xb8/0xd0
2.118677] r7:00000001 r6:8044827c r5:90a99e84 r4:00000000
2.124388] bus_for_each_drv from __device_attach+0xe0/0x158
2.130226] r6:82173054 r5:8137c268 r4:82173010
2.134881] __device_attach from device_initial_probe+0x1c/0x20
2.140977] r7:8200f000 r6:82173010 r5:8137c268 r4:82173010
2.146689] device_initial_probe from bus_probe_device+0x38/0x90
2.152869] bus_probe_device from deferred_probe_work_func+0x90/0xa8
2.159400] r7:8200f000 r6:8137c1ec r5:8137c188 r4:82173010
2.165121] deferred_probe_work_func from process_one_work+0x1ac/0x278
2.171839] r7:8200f000 r6:82006800 r5:8137c1b0 r4:82180000
2.177550] process_one_work from process_scheduled_works+0x38/0x3c
2.184011] r10:90841e94 r9:81399b20 r8:8200681c r7:80d03d00 r6:82006800 r5:82180018
2.191913] r4:82180000
2.194478] process_scheduled_works from worker_thread+0x23c/0x2d4
2.200840] r5:82180018 r4:82180000
2.204455] worker_thread from kthread+0xd4/0xdc
2.209254] r9:82180000 r8:80131078 r7:821e2150 r6:821e2080 r5:821e2140 r4:8219cac0
2.217865] kthread from ret_from_fork+0x14/0x2c
```

Backtrace

```
2.221842] Exception stack(0x90a99fb0 to 0x90a99ff8)
2.226958] 9fa0: 00000000 00000000 00000000 00000000 00000000 00000000
2.235220] 9fc0: 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000
2.243479] 9fe0: 00000000 00000000 00000000 00000000 00000013 00000000
2.250170] r10:00000000 r9:00000000 r8:00000000 r7:00000000 r6:00000000 r5:801373c4
2.258074] r4:821e2140 r3:00000001
2.261704] Code: e3500000 0affffe6 e3a03000 e1a00004 (e5936050)
2.267944] ---[ end trace 0000000000000000 ]---
```

Exception stack
content



Kernel oops debugging: `addr2line`

- ▶ In order to convert addresses/symbol name from this display to source code lines, one can use `addr2line`
 - `addr2line -e vmlinux <address>`
- ▶ GNU binutils ≥ 2.39 takes the `symbol+offset` notation too:
 - `addr2line -e vmlinux <symbol_name>+<off>`
- ▶ The `symbol+offset` notation can be used with older binutils versions via the `faddr2line` script in the kernel sources:
 - `scripts/faddr2line vmlinux <symbol_name>+<off>`
- ▶ The kernel must have been compiled with `CONFIG_DEBUG_INFO=y` to embed the debugging information into the `vmlinux` file.



Kernel oops debugging: `decode_stacktrace.sh`

- ▶ `addr2line` decoding of oopses can be automated using `decode_stacktrace.sh` script which is provided in the kernel sources.
- ▶ This script will translate all symbol names/addresses to the matching file/lines and will display the assembly code where the crash did trigger.
- ▶ `./scripts/decode_stacktrace.sh vmlinux [linux_source_path/] \
< oops_report.txt > decoded_oops.txt`
- ▶ NOTE: `CROSS_COMPILE` and `ARCH` env var should be set to obtain the correct disassembly dump.



Panic and oops behavior configuration

- ▶ Sometimes, crash might be so bad that the kernel will panic and halt its execution entirely by stopping scheduling application and staying in a busy loop.
- ▶ Automatic reboot on panic can be enabled via `CONFIG_PANIC_TIMEOUT`
 - 0: never reboots
 - Negative value: reboot immediately
 - Positive value: seconds to wait before rebooting
- ▶ OOPS can be configured to always panic:
 - at boot time, adding `oops=panic` to the command line
 - at build time, setting `CONFIG_PANIC_ON_OOPS=y`



Built-in kernel self tests



Kernel memory issue debugging

- ▶ The same kind of memory issues that can happen in user space can be triggered while writing kernel code
 - Out of bounds accesses
 - Use-after-free errors (dereferencing a pointer after `kfree()`)
 - Out of memory due to missing `kfree()`
- ▶ Various tools are present in the kernel to catch these issues
 - *KASAN* to find use-after-free and out-of-bound memory accesses
 - *KFENCE* to find use-after-free and out-of-bound in production systems
 - *Kmemleak* to find memory leak due to missing free of memory



- ▶ Kernel Address Space Sanitizer
- ▶ Allows to find use-after-free and out-of-bounds memory accesses
- ▶ Uses GCC to instrument the kernel at compile-time
- ▶ Supported by almost all architectures (ARM, ARM64, PowerPC, RISC-V, S390, Xtensa and X86)
- ▶ Needs to be enabled at kernel configuration with `CONFIG_KASAN`
- ▶ Can then be enabled for files by modifying Makefile
 - `KASAN_SANITIZE_file.o := y` for a specific file
 - `KASAN_SANITIZE := y` for all files in the Makefile folder



Kmemleak

- ▶ Kmemleak allows to find memory leaks for dynamically allocated objects with `kmalloc()`
 - Works by scanning the memory to detect if allocated address are not referenced anymore anywhere (large overhead).
- ▶ Once enabled with [CONFIG_DEBUG_KMEMLEAK](#), kmemleak control files will be visible in *debugfs*
- ▶ Memory leaks is scanned every 10 minutes
 - can be disabled via [CONFIG_DEBUG_KMEMLEAK_AUTO_SCAN](#)
- ▶ An immediate scan can be triggered using
 - `# echo scan > /sys/kernel/debug/kmemleak`
- ▶ Results are displayed in *debugfs*
 - `# cat /sys/kernel/debug/kmemleak`
- ▶ See [dev-tools/kmemleak](#) for more information



Kmemleak report

```
# cat /sys/kernel/debug/kmemleak
unreferenced object 0x82d43100 (size 64):
  comm "insmod", pid 140, jiffies 4294943424 (age 270.420s)
  hex dump (first 32 bytes):
    b4 bb e1 8f c8 a4 e1 8f 8c ce e1 8f 88 c6 e1 8f  .....
    10 a5 e1 8f 18 e2 e1 8f ac c6 e1 8f 0c c1 e1 8f  .....
  backtrace:
    [] slab_post_alloc_hook+0xa8/0x1b8
    [] kmem_cache_alloc_trace+0xb8/0x104
    [<1836406b>] 0x7f005038
    [<89fff56d>] do_one_initcall+0x80/0x1a8
    [<31d908e3>] do_init_module+0x50/0x210
    [<2658dd55>] load_module+0x208c/0x211c
    [<e1d48f15>] sys_finit_module+0xe4/0xf4
    [<1de12529>] ret_fast_syscall+0x0/0x54
    [<7ee81f34>] 0x7eca8c80
```



- ▶ UBSAN is a runtime checker for code with undefined behavior
 - Shifting with a value larger than the type
 - Overflow of integers (signed and unsigned)
 - Misaligned pointer access
 - Out of bound access to static arrays
 - <https://clang.llvm.org/docs/UndefinedBehaviorSanitizer.html>
- ▶ It uses compile-time instrumentation to insert checks that will be executed at runtime
- ▶ Must be enabled using `CONFIG_UBSAN=y`
- ▶ Then, can be enabled for specific files by modifying Makefile
 - `UBSAN_SANITIZE_file.o := y` for a specific file
 - `UBSAN_SANITIZE := y` for all files in the Makefile folder



UBSAN: report example

- Report for an undefined behavior due to a shift with a value > 32 .

```
UBSAN: Undefined behaviour in mm/page_alloc.c:3117:19
shift exponent 51 is too large for 32-bit type 'int'
CPU: 0 PID: 6520 Comm: syz-executor1 Not tainted 4.19.0-rc2 #1
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS Bochs 01/01/2011
Call Trace:
__dump_stack lib/dump_stack.c:77 [inline]
dump_stack+0xd2/0x148 lib/dump_stack.c:113
ubsan_epilogue+0x12/0x94 lib/ubsan.c:159
__ubsan_handle_shift_out_of_bounds+0x2b6/0x30b lib/ubsan.c:425
...
RIP: 0033:0x4497b9
Code: e8 8c 9f 02 00 48 83 c4 18 c3 0f 1f 80 00 00 00 00 48 89 f8 48
89 f7 48 89 d6 48 89 ca 4d 89 c2 4d 89 c8 4c 8b 4c 24 08 0f 05 <48> 3d
01 f0 ff ff 0f 83 9b 6b fc ff c3 66 2e 0f 1f 84 00 00 00 00
RSP: 002b:00007fb5ef0e2c68 EFLAGS: 00000246 ORIG_RAX: 0000000000000010
RAX: ffffffffda RBX: 00007fb5ef0e36cc RCX: 00000000004497b9
RDX: 0000000020000040 RSI: 0000000000000258 RDI: 0000000000000014
RBP: 000000000071bea0 R08: 0000000000000000 R09: 0000000000000000
R10: 0000000000000000 R11: 0000000000000246 R12: 00000000ffffffff
R13: 0000000000005490 R14: 00000000006ed530 R15: 00007fb5ef0e3700
```



Debugging locking

- ▶ Lock debugging: prove locking correctness
 - [CONFIG_PROVE_LOCKING](#)
 - Adds instrumentation to kernel locking code
 - Detect violations of locking rules during system life, such as:
 - Locks acquired in different order (keeps track of locking sequences and compares them).
 - Spinlocks acquired in interrupt handlers and also in process context when interrupts are enabled.
 - Not suitable for production systems but acceptable overhead in development.
 - See [locking/lockdep-design](#) for details
- ▶ [CONFIG_DEBUG_ATOMIC_SLEEP](#) allows to detect code that incorrectly sleeps in atomic section (while holding lock typically).
 - Warning displayed in `dmesg` in case of such violation.



Concurrency issues

- ▶ Kernel Concurrency SANitizer framework
- ▶ `CONFIG_KCSAN`, introduced in Linux 5.8.
- ▶ Dynamic race detector relying on compile time instrumentation.
- ▶ Can find concurrency issues (mainly data races) in your system.
- ▶ See dev-tools/kcsan and <https://lwn.net/Articles/816850/> for details.



Debugging kernel programming mistakes with integrated frameworks

- ▶ Debug locking issues using lockdep
- ▶ Spot function calls in invalid context
- ▶ Use kmemleak to detect memory leaks on the system



The Magic SysRq



The Magic SysRq

Functionality provided by serial drivers

- ▶ Allows to run multiple debug/rescue commands even when the kernel seems to be in deep trouble
 - On embedded: in the console, send a break character (Picocon: press [Ctrl] + a followed by [Ctrl] + \), then press <character>
 - By echoing <character> in /proc/sysrq-trigger
- ▶ Example commands:
 - h: show available commands
 - s: sync all mounted filesystems
 - b: reboot the system
 - w: shows the kernel stack of all sleeping processes
 - t: shows the kernel stack of all running processes
 - g: enter kgdb mode
 - z: flush trace buffer
 - c: triggers a crash (kernel panic)
 - You can even register your own!
- ▶ Detailed in [admin-guide/sysrq](#)



KGDB



kgdb - A kernel debugger

- ▶ `CONFIG_KGDB` in *Kernel hacking*.
- ▶ The execution of the kernel is fully controlled by `gdb` from another machine, connected through a serial line.
- ▶ Can do almost everything, including inserting breakpoints in interrupt handlers.
- ▶ Feature supported for the most popular CPU architectures
- ▶ `CONFIG_GDB_SCRIPTS` allows to build GDB python scripts that are provided by the kernel.
 - See [process/debugging/kgdb](https://bootlin.com/process/debugging/kgdb) for more information



kgdb kernel config

- ▶ `CONFIG_DEBUG_KERNEL=y` to make KGDB support visible
- ▶ `CONFIG_KGDB=y` to enable KGDB support
- ▶ `CONFIG_DEBUG_INFO=y` to compile the kernel with debug info (`-g`)
- ▶ `CONFIG_FRAME_POINTER=y` to have more reliable stacktraces
- ▶ `CONFIG_KGDB_SERIAL_CONSOLE=y` to enable KGDB support over serial
- ▶ `CONFIG_GDB_SCRIPTS=y` to enable kernel GDB python scripts
- ▶ `CONFIG_RANDOMIZE_BASE=n` to disable KASLR
- ▶ `CONFIG_WATCHDOG=n` to disable watchdog
- ▶ `CONFIG_MAGIC_SYSRQ=y` to enable Magic SysReq support
- ▶ `CONFIG_STRICT_KERNEL_RWX=n` to disable memory protection on code section, thus allowing to put breakpoints



kgdb pitfalls

- ▶ KASLR should be disabled to avoid confusing gdb with randomized kernel addresses
 - Disable *kaslr mode* using `nokaslr` command line parameter if enabled in your kernel.
- ▶ Disable the platform watchdog to avoid rebooting while debugging.
 - When interrupted by KGDB, all interrupts are disabled thus, the watchdog is not serviced.
 - Sometimes, watchdog is enabled by upper boot levels. Make sure to disable the watchdog there too.
- ▶ Can not interrupt kernel execution from gdb using `interrupt` command or `Ctrl + C`.
- ▶ Not possible to break everywhere (see [CONFIG_KGDB_HONOUR_BLOCKLIST](#)).
- ▶ Need a console driver with polling support.
- ▶ Some architecture lacks functionalities (No watchpoints on arm32 for instance) and some instabilities might happen!



Using kgdb (1/2)

- ▶ Details available in the kernel documentation: [process/debugging/kgdb](#)
- ▶ You must include a kgdb I/O driver. One of them is kgdb over serial console (kgdboc: kgdb over console, enabled by [CONFIG_KGDB_SERIAL_CONSOLE](#))
- ▶ Configure kgdboc at boot time by passing to the kernel:
 - `kgdboc=<tty-device>,<bauds>.`
 - For example: `kgdboc=ttys0,115200`
- ▶ Or at runtime using sysfs:
 - `echo ttys0 > /sys/module/kgdboc/parameters/kgdboc`
 - If the console does not have polling support, this command will yield an error.



Using kgdb (2/2)

- ▶ Then also pass `kgdbwait` to the kernel: it makes `kgdb` wait for a debugger connection.
- ▶ Boot your kernel, and when the console is initialized, interrupt the kernel with a break character and then `g` in the serial console (see our *Magic SysRq* explanations).
- ▶ On your workstation, start `gdb` as follows:
 - `arm-linux-gdb ./vmlinux`
 - `(gdb) set remotebaud 115200`
 - `(gdb) target remote /dev/ttyS0`
- ▶ Once connected, you can debug a kernel the way you would debug an application program.
- ▶ On GDB side, the first threads represent the CPU context (`ShadowCPU<x>`), then all the other threads represents a task.



Kernel *GDB* scripts

- ▶ `CONFIG_GDB_SCRIPTS` allows to build a set of python script which ease the kernel debugging by adding new commands and functions.
- ▶ When using `gdb vmlinux`, the scripts present in `vmlinux-gdb.py` file at the root of build dir will be loaded automatically.
 - `lx-symbols`: (Re)load symbols for `vmlinux` and modules
 - `lx-dmesg`: display kernel dmesg
 - `lx-lsmod`: display loaded modules
 - `lx-device-{bus|class|tree}`: display device bus, classes and tree
 - `lx-ps`: ps like view of tasks
 - `$lx_current()` contains the current `task_struct`
 - `$lx_per_cpu(var, cpu)` returns a per-cpu variable
 - apropos `lx` To display all available functions.
- ▶ [dev-tools/gdb-kernel-debugging](#)



- ▶ `CONFIG_KGDB_KDB` includes a kgdb frontend name "KDB"
- ▶ This frontend exposes a debug prompt on the serial console which allows debugging the kernel without the need for an external gdb.
- ▶ KDB can be entered using the same mechanism used for entering kgdb mode.
- ▶ *KDB* and *KGDB* can coexist and be used at the same time.
 - Use the `kgdb` command in KDB to enter kgdb mode.
 - Send a maintenance packet from gdb using `maintenance packet 3` to switch from kgdb to KDB mode.



KDB commands


- ▶ KDB does not consume gdb commands but a set of dedicated KDB commands:
 - go: Continue execution
 - bt: Display backtrace
 - env: Show environment variables
 - ps: List all tasks
 - pid: Switch to another task
 - md/mm: Read/write memory
 - lsmod: List loaded modules
- ▶ To check all available commands, you can refer to the `help` command output, or check `maintab` in kernel source code



- ▶ When the system has only a single serial port, it is not possible to use both KGDB and the serial line as an output terminal since only one program can access that port.
- ▶ Fortunately, the *kdmx* tool allows to use both KGDB and serial output by splitting GDB messages and standard console from a single port to 2 slave pty (/dev/pts/x)
- ▶ <https://git.kernel.org/pub/scm/utils/kernel/kgdb/agent-proxy.git>
 - Located in the subdirectory *kdmx*

```
$ kdmx -n -d -p/dev/ttyACM0 -b115200
serial port: /dev/ttyACM0
Initializing the serial port to 115200 8n1
/dev/pts/6 is slave pty for terminal emulator
/dev/pts/7 is slave pty for gdb


Use <ctrl>C to terminate program
```





Going further with KGDB

- ▶ Good presentation from Doug Anderson with a lot of demos and explanations
 - Video: <https://www.youtube.com/watch?v=HB0woSyRmys>
 - Slides: https://elinux.org/images/1/1b/ELC19_Serial_kdb_kgdb.pdf





crash



- ▶ *crash* is a CLI tool allowing to investigate kernel (dead or alive!)
 - Uses `/dev/mem` or `/proc/kcore` on live systems
 - Requires `CONFIG_STRICT_DEVMEM=n`
- ▶ Can use a coredump generated using `kdump`, `kvmdump`, etc.
- ▶ Based on `gdb` and provides many specific commands to inspect the kernel state.
 - Stack traces, `dmesg` (log), memory maps of the processes, `irqs`, virtual memory areas, etc.
- ▶ Allows examining all the tasks that are running on the system.
- ▶ Hosted at <https://github.com/crash-utility/crash>



crash example

```
$ crash vmlinux vmcore
[...]  
  TASKS: 75  
NODENAME: buildroot  
RELEASE: 5.13.0  
VERSION: #1 SMP PREEMPT Tue Nov 15 14:42:25 CET 2022  
MACHINE: armv7l (unknown Mhz)  
MEMORY: 512 MB  
  PANIC: "Unable to handle kernel NULL pointer dereference at virtual address 00000070"  
    PID: 127  
  COMMAND: "watchdog"  
    TASK: c3f163c0 [THREAD_INFO: c3f00000]  
    CPU: 1  
    STATE: TASK_RUNNING (PANIC)  
  
crash> mach  
  MACHINE TYPE: armv7l  
  MEMORY SIZE: 512 MB  
    CPUS: 1  
  PROCESSOR SPEED: (unknown)  
    HZ: 100  
  PAGE SIZE: 4096  
  KERNEL VIRTUAL BASE: c0000000  
  KERNEL MODULES BASE: bf000000  
  KERNEL VMALLOC BASE: e0000000  
  KERNEL STACK SIZE: 8192
```




Debugging kernel crashes on a live kernel

- ▶ Analyze an OOPS message
- ▶ Debug a crash with KGDB



Post-mortem analysis



Kernel crash post-mortem analysis

- ▶ Sometimes, accessing the crashed system is not possible or the system can't stay offline while waiting to be debugged
- ▶ Kernel can generate crash dumps (a *vmcore* file) to a remote location, allowing to quickly restart the system while still be able to perform post-mortem analysis with GDB.
- ▶ This feature relies on *kexec* and *kdump* which will boot another kernel as soon as the crash occurs right after dumping the *vmcore* file.
 - The *vmcore* file can be saved on local storage, via SSH, FTP etc.



kexec & kdump (1/2)

- ▶ On panic, the kernel kexec support will execute a "dump-capture kernel" directly from the kernel that crashed
 - Most of the time, a specific dump-capture kernel is compiled for that task (minimal config with specific initramfs/initrd)
- ▶ *kexec* system works by saving some RAM for the kdump kernel execution at startup
 - `crashkernel` parameter should be set to specify the crash kernel dedicated physical memory region
- ▶ *kexec-tools* are then used to load dump-capture kernel into this memory zone using the `kexec` command
 - Internally uses the `kexec_load` system call [man 2 kexec_load](#)



kexec & kdump (2/2)

- ▶ Finally, on panic, the kernel will reboot into the "dump-capture" kernel allowing the user to dump the kernel coredump (`/proc/vmcore`) onto whatever media
- ▶ Additional command line options depends on the architecture
- ▶ See [admin-guide/kdump/kdump](#) for more comprehensive explanations on how to setup the kdump kernel with `kexec`.
- ▶ Additional user-space services and tools allow to automatically collect and dump the vmcore file to a remote location.
 - See kdump systemd service and the `makedumpfile` tool which can also compress the vmcore file into a smaller file (Only for x86, PPC, IA64, S390).
 - <https://github.com/makedumpfile/makedumpfile>



kdump




Image credits: Wikipedia




kexec config and setup

- ▶ On the standard kernel:
 - `CONFIG_KEXEC=y` to enable KEXEC support
 - `kexec-tools` to provide the `kexec` command
 - A kernel and a DTB accessible by `kexec`
- ▶ On the dump-capture kernel:
 - `CONFIG_CRASH_DUMP=y` to enable dumping a crashed kernel
 - `CONFIG_PROC_VMCORE=y` to enable `/proc/vmcore` support
 - `CONFIG_AUTO_ZRELADDR=y` on ARM32 platforms
- ▶ Set the correct `crashkernel` command line option:
 - `crashkernel=size[KMG][@offset[KMG]]`
- ▶ Load a dump-capture kernel on the first kernel with `kexec`:
 - `kexec --type zImage -p my_zImage --dtb=my_dtb.dtb --initrd=my_initrd --append="command line option"`
- ▶ Then simply wait for a crash to happen!



Going further with kexec & kdump

- ▶ Presentation from Steven Rostedt about using kexec, kdump and ftrace with lot of tips and tricks about using kexec/kdump
 - Video: <https://www.youtube.com/watch?v=aUGNDJPpUUg>
 - Slides: https://static.sched.com/hosted_files/ossna2022/c0/Postmortem_%20Kexec%2C%20Kdump%20and%20Ftrace.pdf





pstore (1/2)

- ▶ Linux provides a filesystem interface for Persistent Storage (`pstore`) to save data across system resets: kernel logs, oopses, ftrace records, user messages...
- ▶ The platform needs to provide a persistent area to `pstore` (a block device, reserved RAM which is not reset on reboot, etc). Then you can enable a `pstore` frontend.
- ▶ `ramoops` is a common frontend for `pstore`: it will log any panic/oops to a `pstore`-managed ram buffer, which will be accessible on next boot
- ▶ Saved logs can be retrieved on next boot thanks to the `pstore` filesystem
- ▶ Some earlier software components in the boot chain (eg: `U-Boot`), if properly configured, may be able to access `pstore` data as well



pstore (2/2)

- ▶ Kernel configuration:
 - `CONFIG_PSTORE=y`
 - `CONFIG_PSTORE_RAM=y`
- ▶ Platform configuration: reserve some memory for pstore and configure it
 - Either through kernel command line:
`mem=<usable_memory_size> ramoops.mem_address=0x80000000 ramoops.ecc=1`
 - Or through device tree:

```
reserved-memory {  
    [...]  
    ramoops@8f000000 {  
        compatible = "ramoops";  
        reg = <0 0x8f000000 0 0x100000>;  
        record-size = <0x4000>;  
        console-size = <0x4000>;  
    };  
};
```

- ▶ After a crash, the collected logs/traces will be available in the pstore filesystem:

```
mount -t pstore pstore /sys/fs/pstore
```




Post-mortem debugging of a kernel crash

- ▶ Setup kexec, kdump and extract a kernel coredump



Going further

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!






Debugging resources

- ▶ Brendan Gregg [Systems performance](#) book
- ▶ Brendan Gregg [Linux Performance](#) page
- ▶ *Tools and Techniques to Debug an Embedded Linux System*, talk from Sergio Prado, [video](#), [slides](#)
- ▶ *Tracing with Ftrace: Critical Tooling for Linux Development*, talk from Steven Rostedt, [video](#)
- ▶ *Tutorial: Debugging Embedded Devices using GDB*, tutorial from Chris Simmonds, [video](#)



Going further (Tracing & Profiling)

- ▶ Great book from Brendan Gregg, an expert in tracing and profiling
- ▶ <https://www.brendangregg.com/blog/2020-07-15/systems-performance-2nd-edition.html>
- ▶ Covers concepts, strategy, tools, and tuning for Linux kernel and applications.





Going further (BPF)

- ▶ Still from Brendan Gregg!
- ▶ Covers more than 150 tools that use BPF.
- ▶ Explains how to analyze the results from these tools to optimize your system.
- ▶ <https://www.brendangregg.com/bpf-performance-tools-book.html>

BPF Performance Tools

Linux System and
Application Observability

Brendan Gregg



Foreword by Alexei Starovoitov,
creator of the new BPF

ADDITION-WESLEY PROFESSIONAL COMPUTING SERIES



Last slides

© Copyright 2004-2025, Bootlin.
Creative Commons BY-SA 3.0 license.
Corrections, suggestions, contributions and translations are welcome!





Thank you!
And may the Source be with you



Rights to copy

© Copyright 2004-2025, Bootlin

License: Creative Commons Attribution - Share Alike 3.0

<https://creativecommons.org/licenses/by-sa/3.0/legalcode>

You are free:

- ▶ to copy, distribute, display, and perform the work
- ▶ to make derivative works
- ▶ to make commercial use of the work

Under the following conditions:

- ▶ **Attribution.** You must give the original author credit.
- ▶ **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
- ▶ For any reuse or distribution, you must make clear to others the license terms of this work.
- ▶ Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

Document sources: <https://github.com/bootlin/training-materials/>